

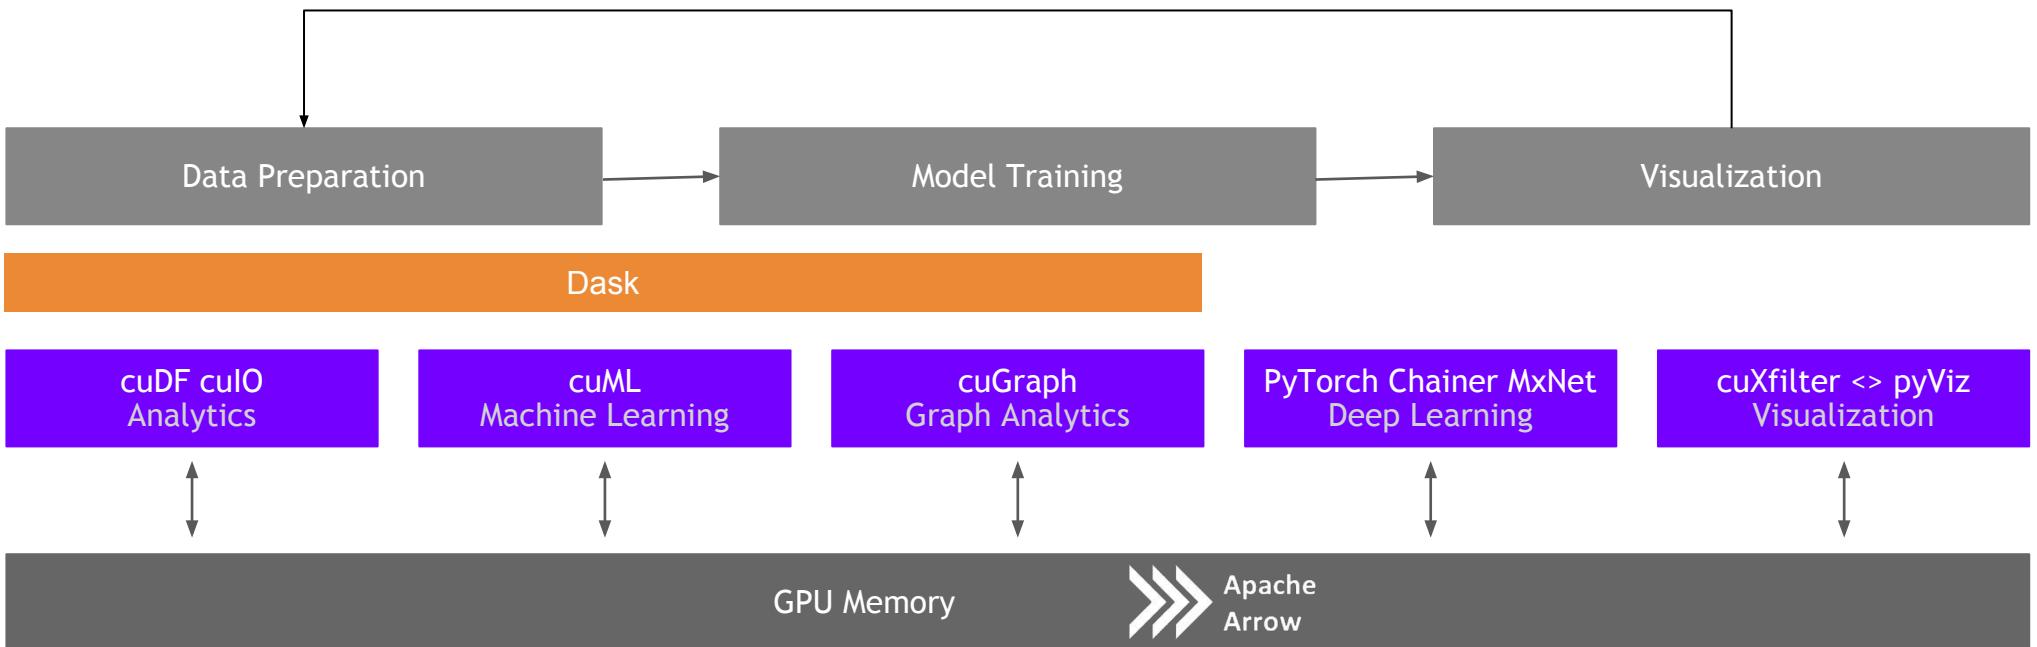
# RAPIDS

The Platform Inside and Out

Joshua Patterson - Director, RAPIDS Engineering

# RAPIDS

## End-to-End Accelerated GPU Data Science



# Data Processing Evolution

Faster data access, less data movement

Hadoop Processing, Reading from disk

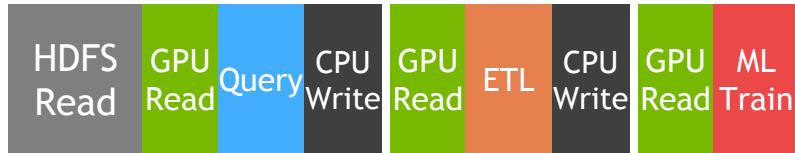


Spark In-Memory Processing



25-100x Improvement  
Less code  
Language flexible  
Primarily In-Memory

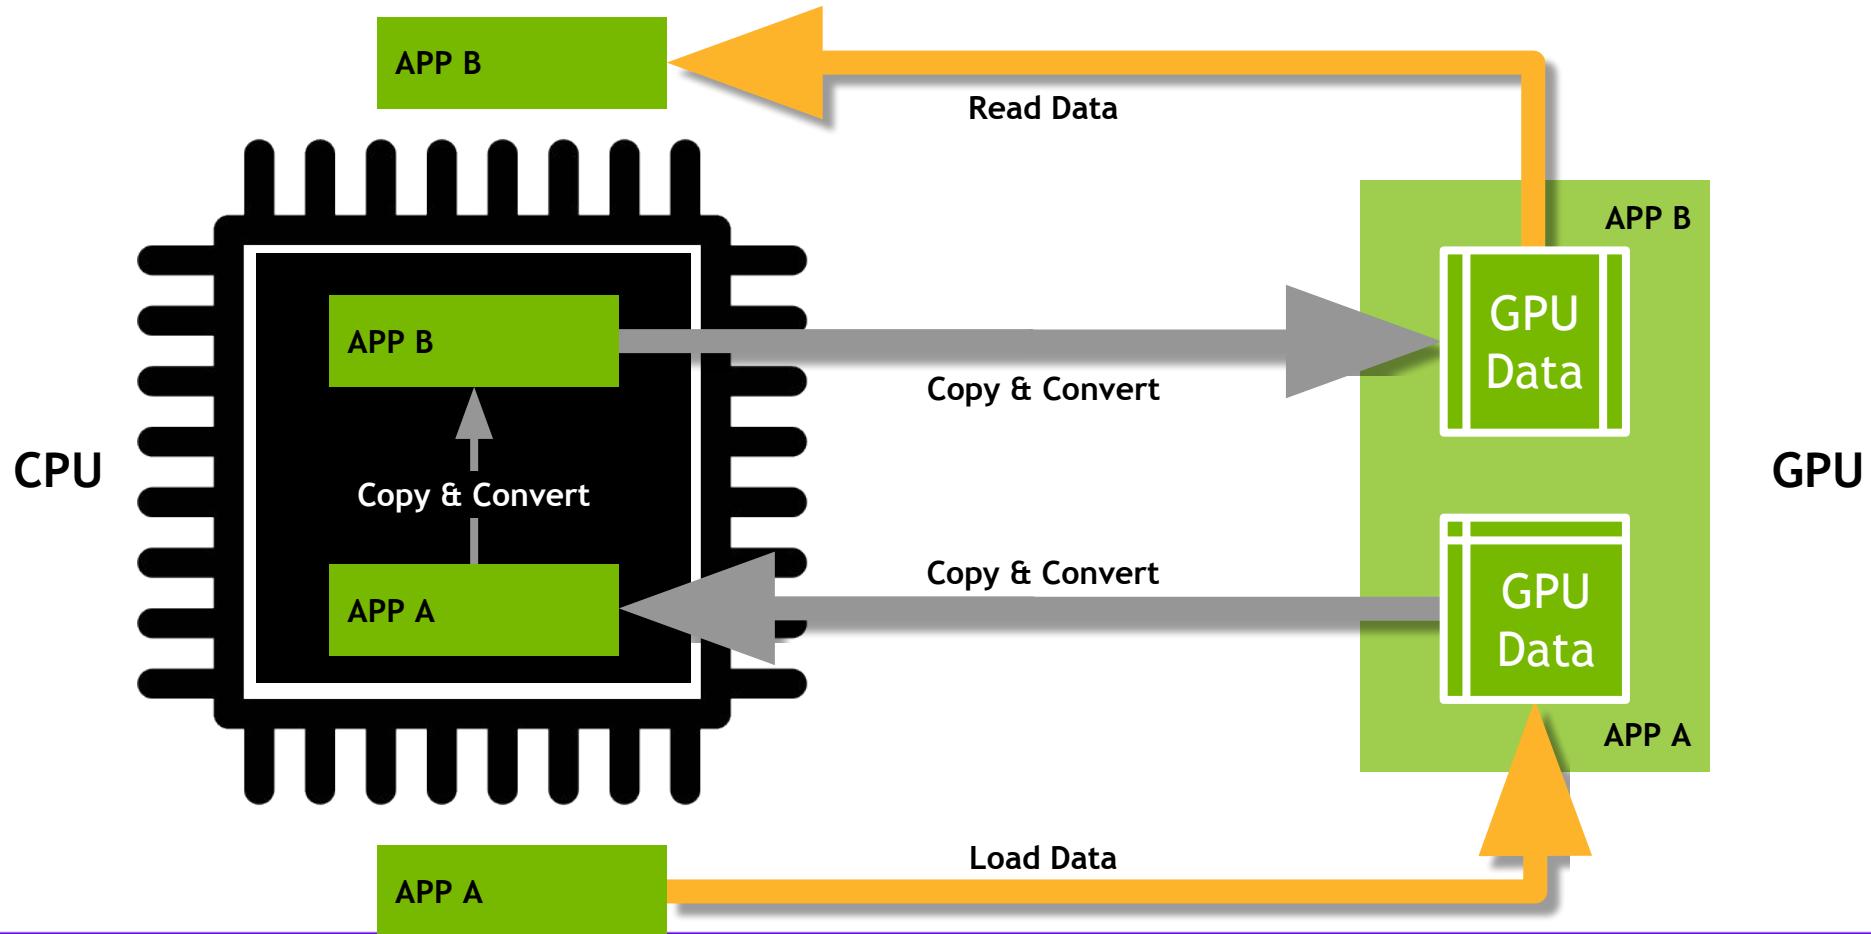
Traditional GPU Processing



5-10x Improvement  
More code  
Language rigid  
Substantially on GPU

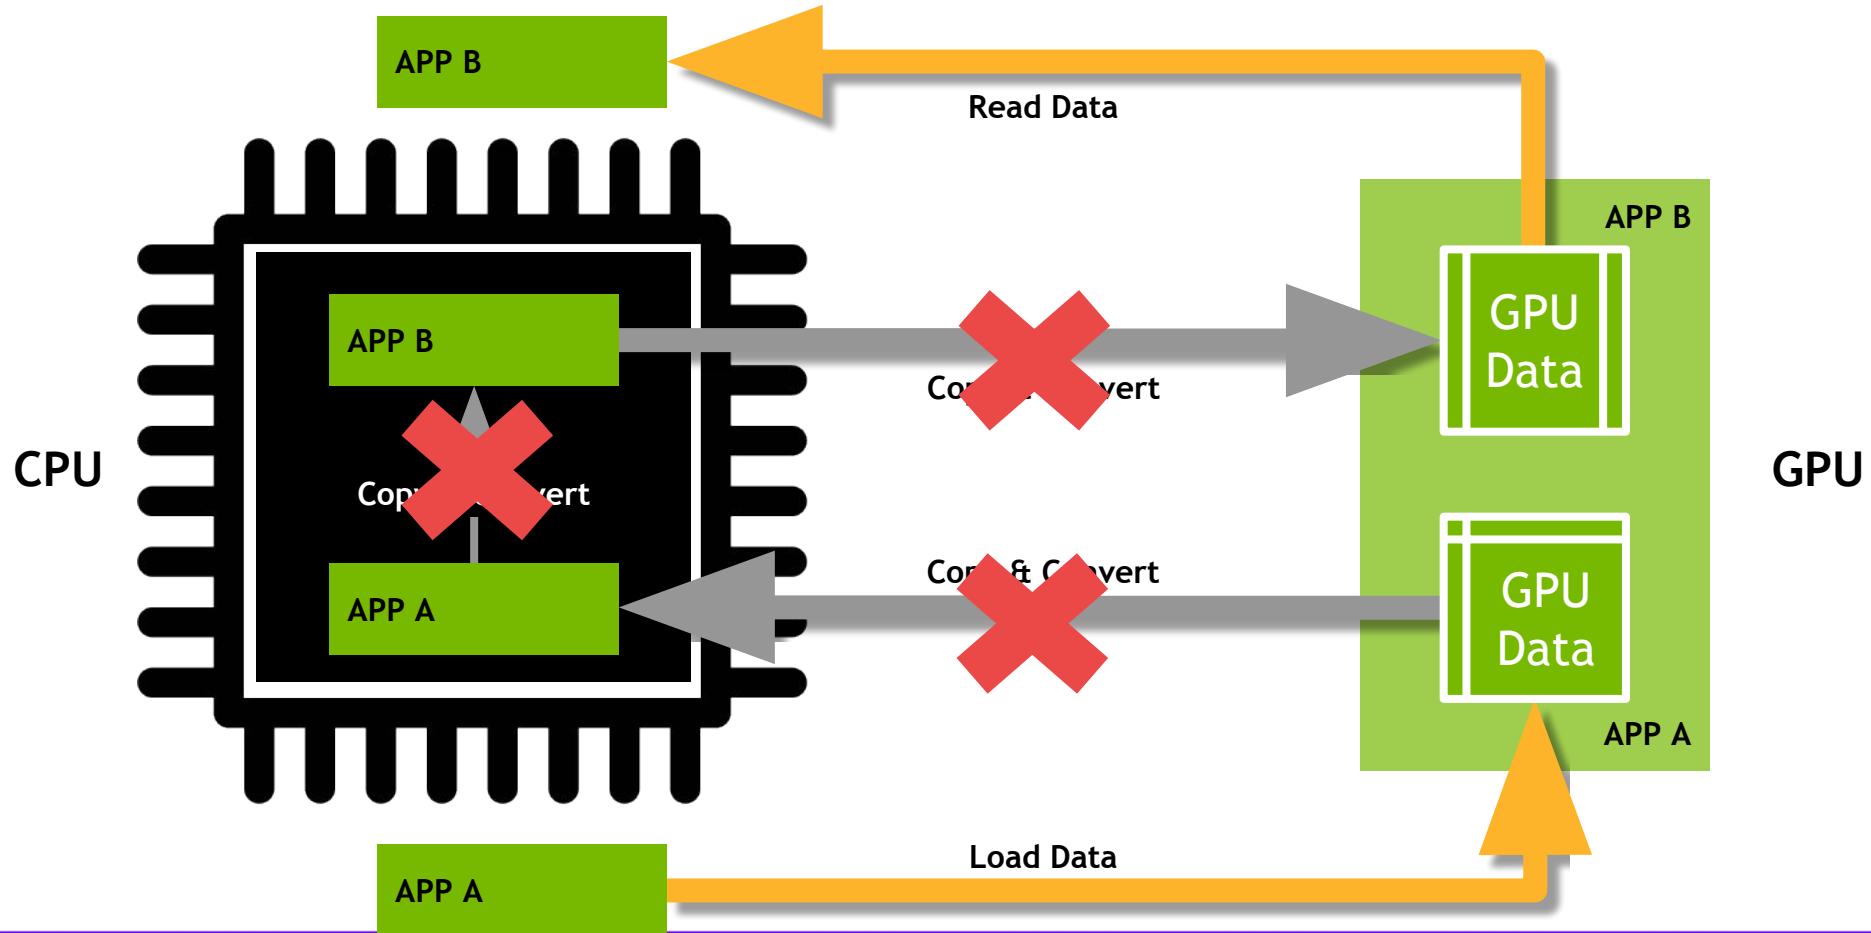
# Data Movement and Transformation

The bane of productivity and performance

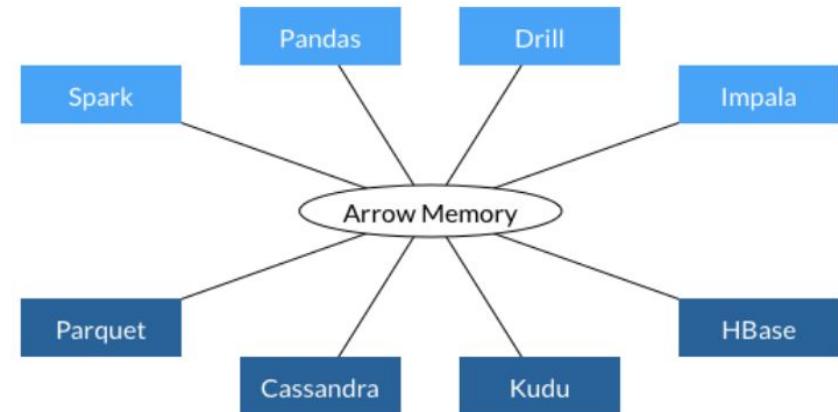
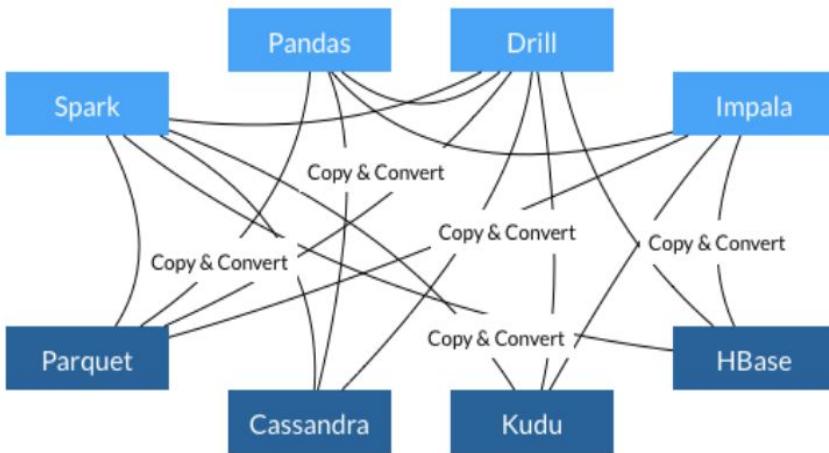


# Data Movement and Transformation

What if we could keep data on the GPU?



# Learning from Apache Arrow ➤➤➤



- Each system has its own internal memory format
- 70-80% computation wasted on serialization and deserialization
- Similar functionality implemented in multiple projects
- All systems utilize the same memory format
- No overhead for cross-system communication
- Projects can share functionality (eg, Parquet-to-Arrow reader)

From Apache Arrow Home Page - <https://arrow.apache.org/>

# Data Processing Evolution

Faster data access, less data movement

Hadoop Processing, Reading from disk

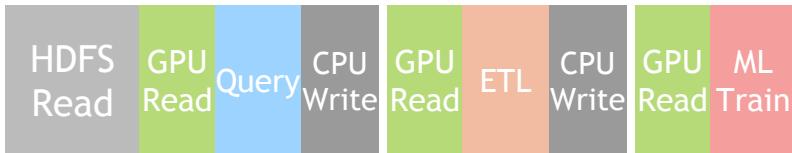


Spark In-Memory Processing



25-100x Improvement  
Less code  
Language flexible  
Primarily In-Memory

Traditional GPU Processing



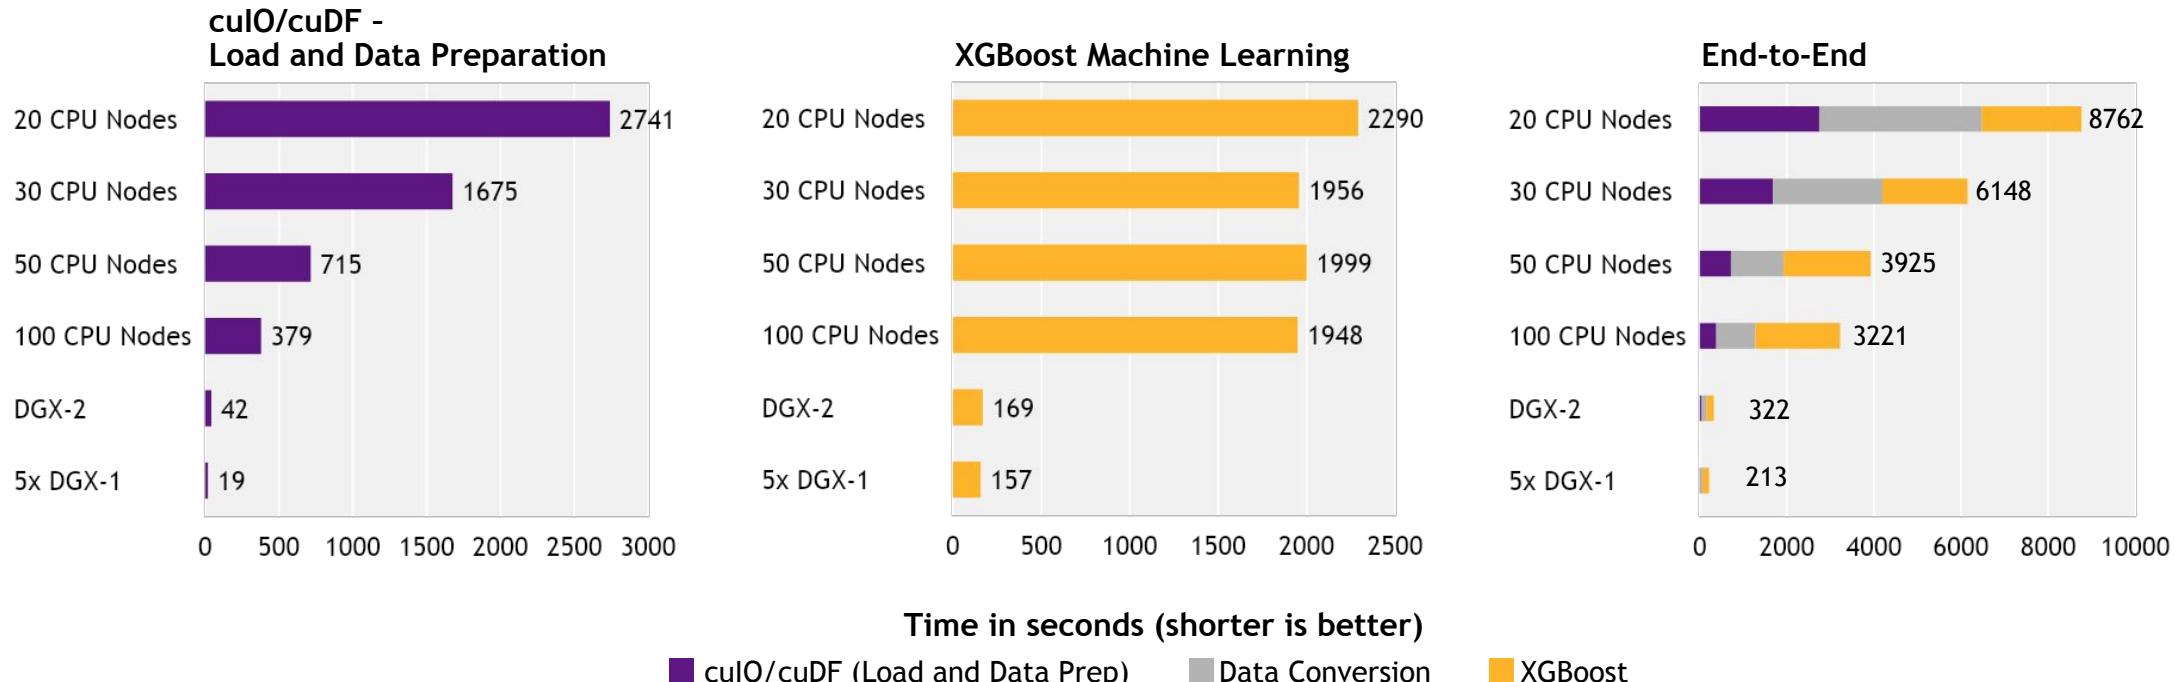
5-10x Improvement  
More code  
Language rigid  
Substantially on GPU

RAPIDS



50-100x Improvement  
Same code  
Language flexible  
Primarily on GPU

# Faster Speeds, Real-World Benefits



## Benchmark

200GB CSV dataset; Data prep includes joins, variable transformations

## CPU Cluster Configuration

CPU nodes (61 GiB memory, 8 vCPUs, 64-bit platform), Apache Spark

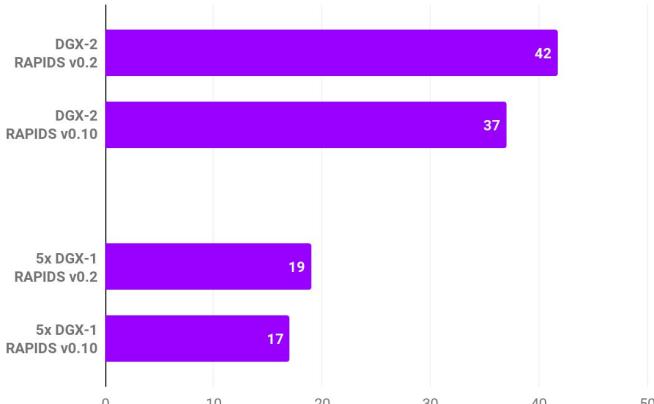
## DGX Cluster Configuration

5x DGX-1 on InfiniBand network

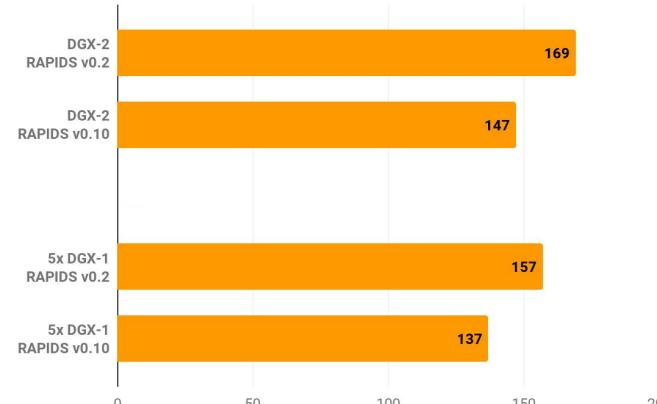
# Faster Speeds, Real-World Benefits

## Improving Over Time

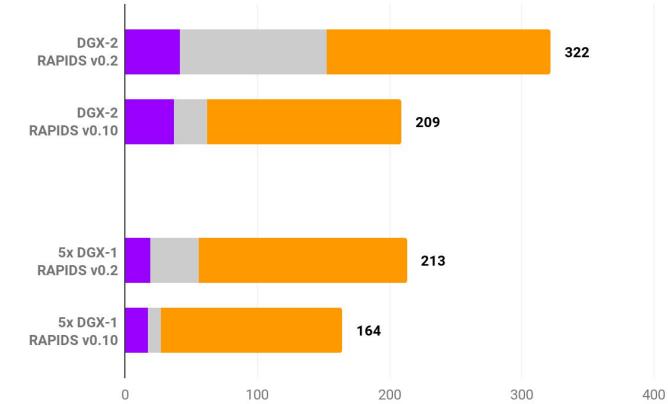
**cuIO/cuDF -  
Load and Data Preparation**



**XGBoost Machine Learning**



**End-to-End**



Time in seconds (shorter is better)

■ cuIO/cuDF (Load and Data Prep)

■ Data Conversion

■ XGBoost

### Benchmark

200GB CSV dataset; Data prep includes joins, variable transformations

### CPU Cluster Configuration

CPU nodes (61 GiB memory, 8 vCPUs, 64-bit platform), Apache Spark

### DGX Cluster Configuration

5x DGX-1 on InfiniBand network

# Speed, UX, and Iteration

## The Way to Win at Data Science

François Chollet @fchollet Following

Winners are those who went through \*more iterations\* of the "loop of progress" -- going from an idea, to its implementation, to actionable results. So the winning teams are simply those able to run through this loop \*faster\*.

And this is where Keras gives you an edge.

12:31 PM - 3 April 2019

50 Retweets 158 Likes

5 50 158

François Chollet @fchollet · Apr 3

We often talk about how following UX best practices for API design makes Keras more accessible and easier to use, and how this helps beginners. But those who stand to benefit most from good UX \*aren't\* the beginners. It's actually the very best practitioners in the world.

1 7 50

François Chollet @fchollet · Apr 3

Because good UX reduces the overhead (development overhead & cognitive overhead) to setting up new experiments. It means you will be able to iterate faster. You will be able to try more ideas.

2 11 78

François Chollet @fchollet · Apr 3

So I don't think it's mere personal preference if Kaggle champions are overwhelmingly using Keras.

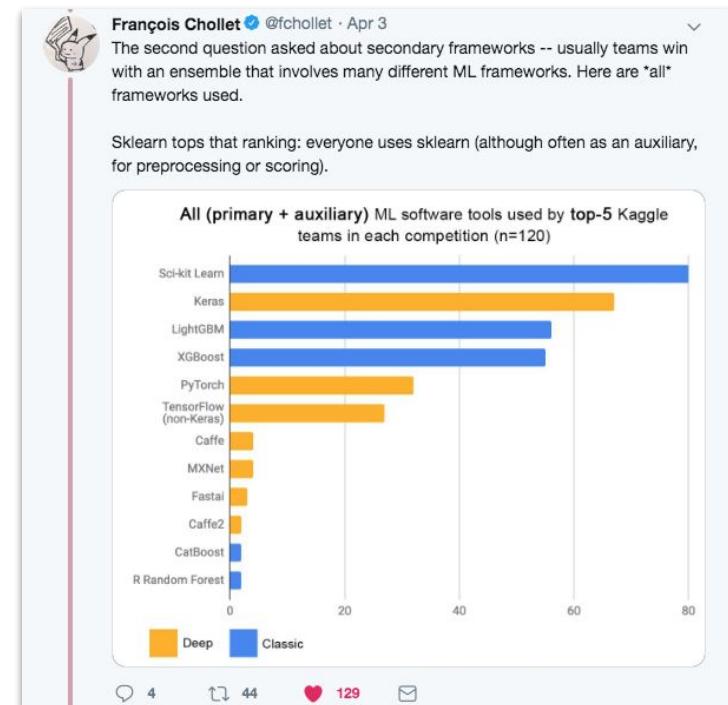
8 8 74

Joshua Patterson @datametrician · Apr 3

Replying to @fchollet

This is the fundamental belief that drives @RAPIDSai. @nvidia #GPU infrastructure is fast, people need to iterate quickly, people want a known #python interface. Combine them and you're off to the races!

2 11

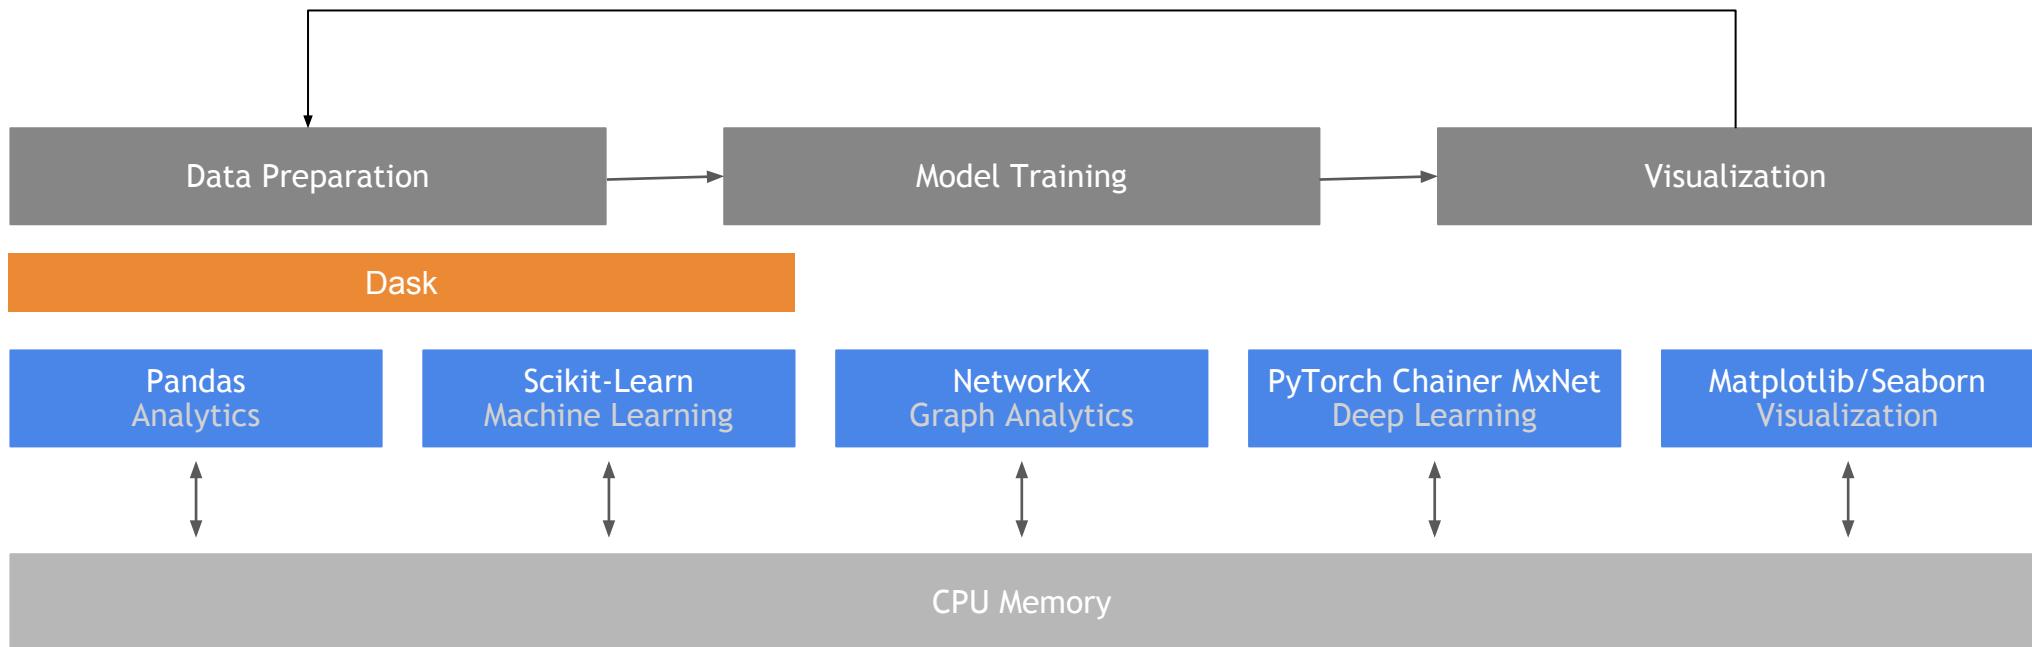


kaggle

# RAPIDS Core

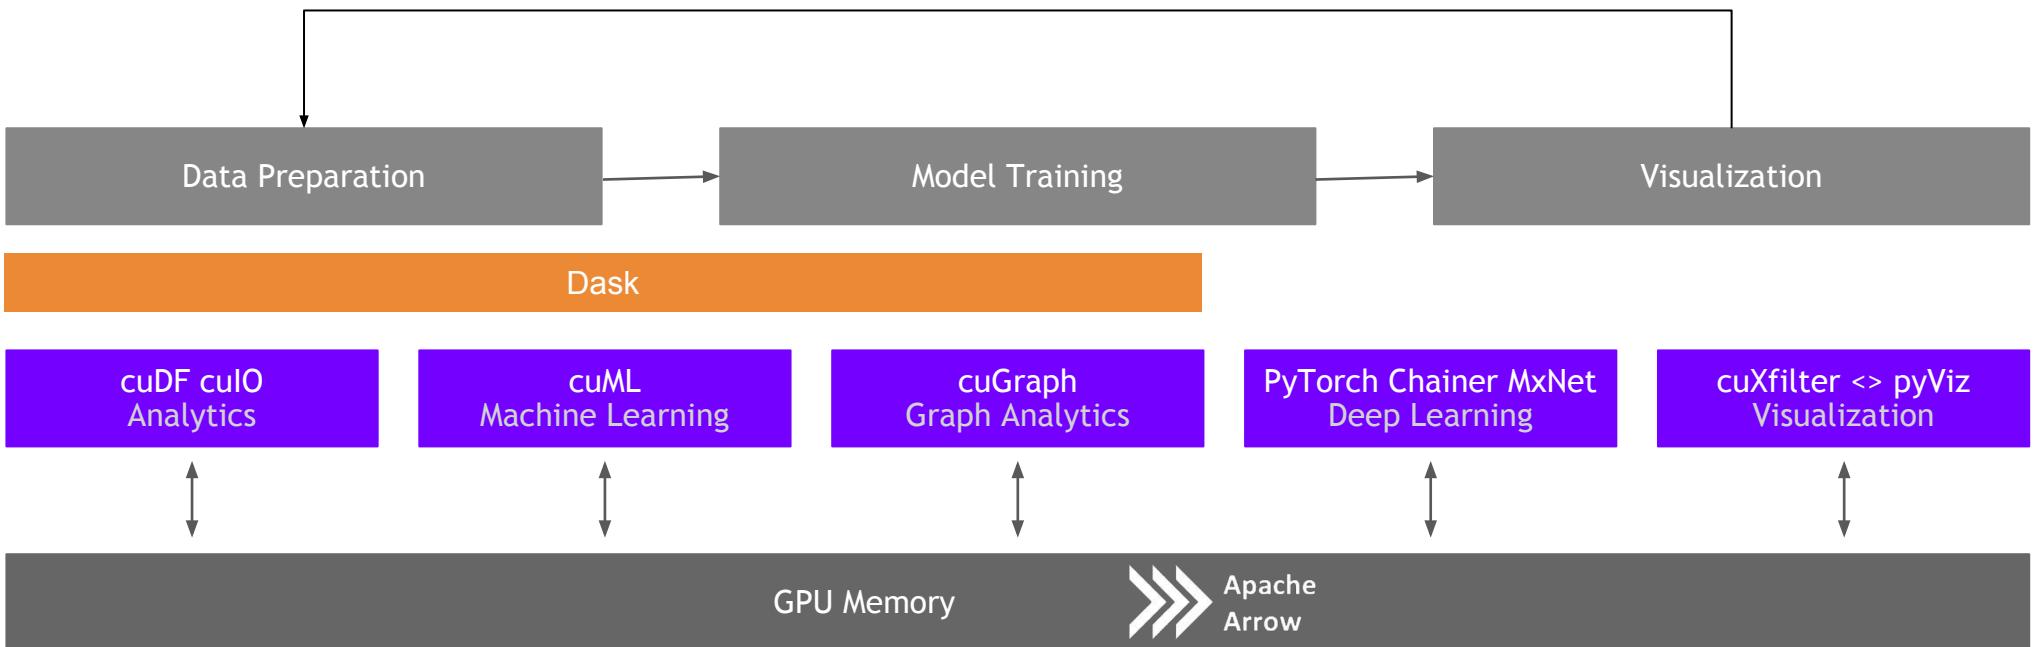
# Open Source Data Science Ecosystem

## Familiar Python APIs



# RAPIDS

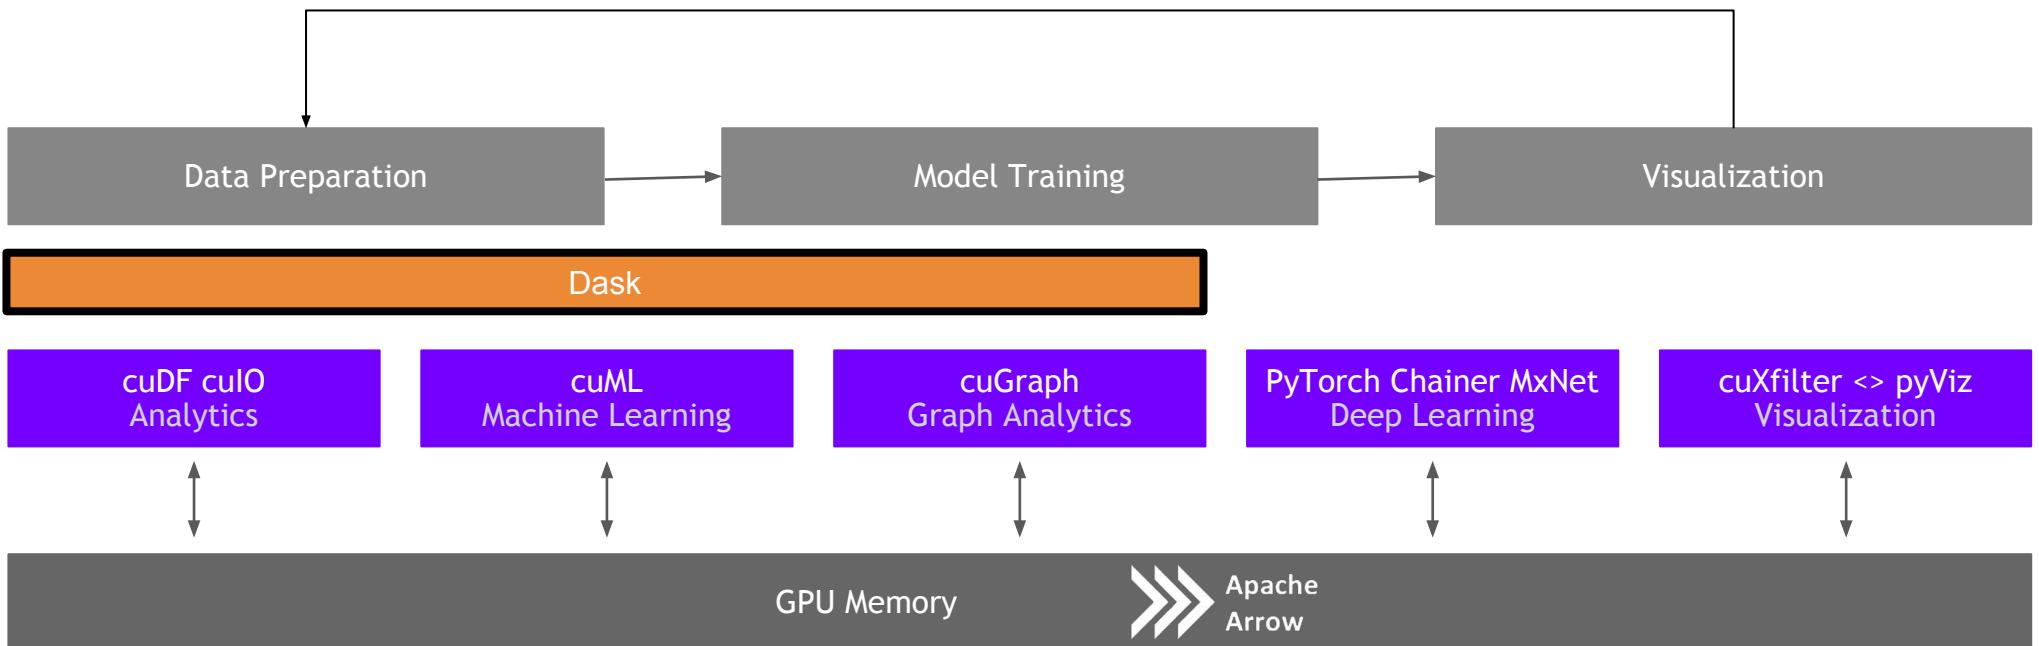
## End-to-End Accelerated GPU Data Science



# Dask

# RAPIDS

## Scaling RAPIDS with Dask



# Why Dask?

## PyData Native

- **Easy Migration:** Built on top of NumPy, Pandas, Scikit-Learn, etc.
- **Easy Training:** With the same APIs
- **Trusted:** With the same developer community

## Deployable

- HPC: SLURM, PBS, LSF, SGE
- Cloud: Kubernetes
- Hadoop/Spark: Yarn



## Easy Scalability

- Easy to install and use on a laptop
- Scales out to thousand-node clusters

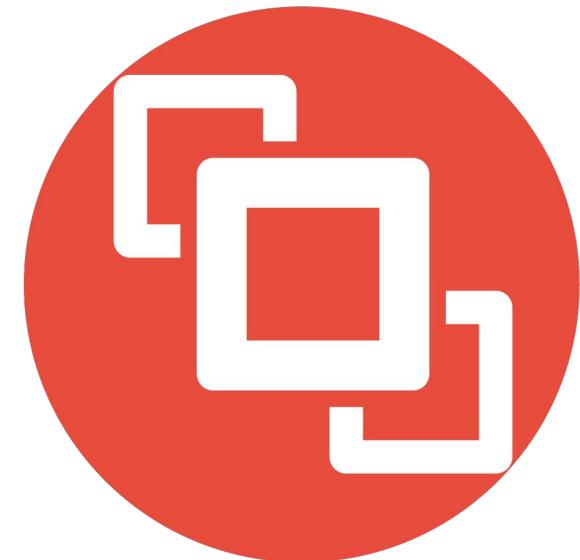
## Popular

- Most common parallelism framework today in the PyData and SciPy community

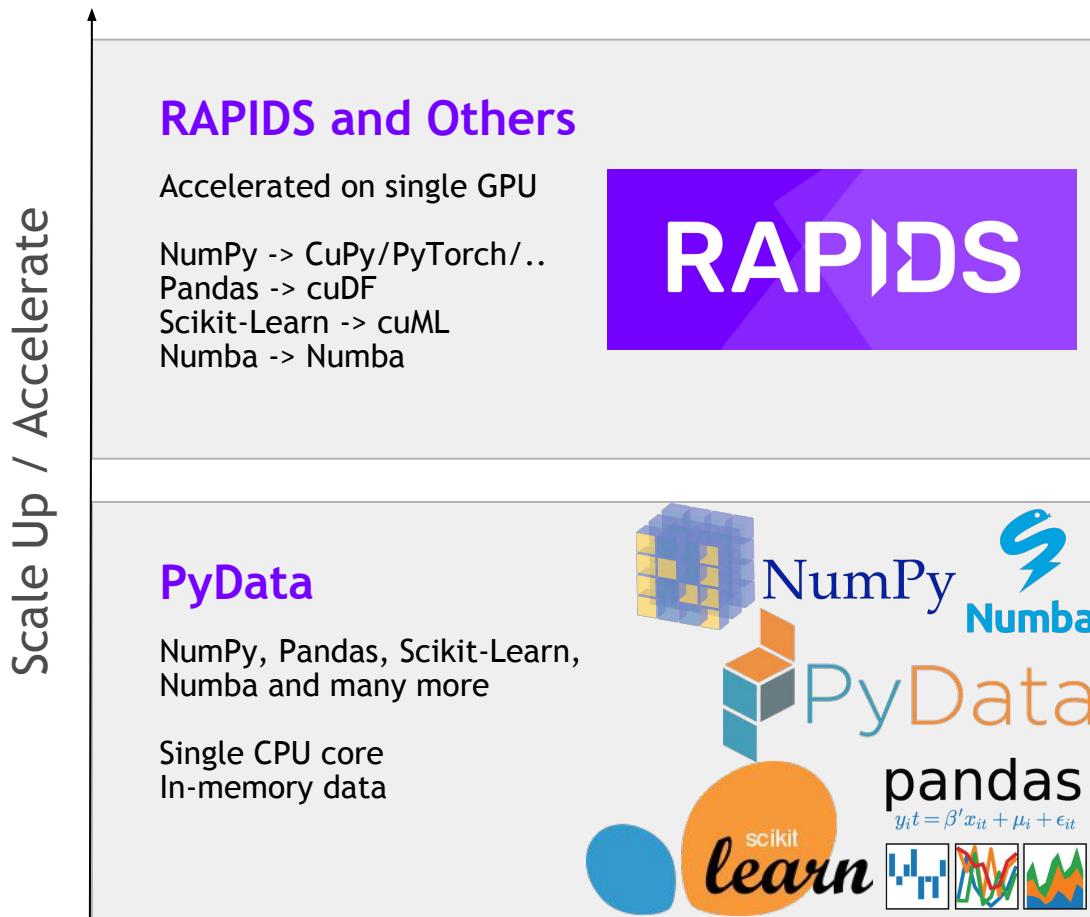
# Why OpenUCX?

Bringing hardware accelerated communications to Dask

- TCP sockets are slow!
- UCX provides uniform access to transports (TCP, InfiniBand, shared memory, NVLink)
- Alpha Python bindings for UCX (ucx-py)  
<https://github.com/rapidsai/ucx-py>
- Will provide best communication performance, to Dask based on available hardware on nodes/cluster



# Scale up with RAPIDS



# Scale out with RAPIDS + Dask with OpenUCX

Scale Up / Accelerate

## RAPIDS and Others

Accelerated on single GPU

NumPy -> CuPy/PyTorch/..  
Pandas -> cuDF  
Scikit-Learn -> cuML  
Numba -> Numba



## RAPIDS + Dask with OpenUCX

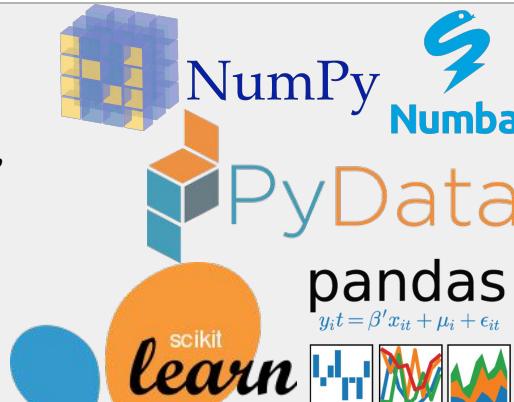
Multi-GPU  
On single Node (DGX)  
Or across a cluster



## PyData

NumPy, Pandas, Scikit-Learn,  
Numba and many more

Single CPU core  
In-memory data



## Dask

Multi-core and Distributed PyData

NumPy -> Dask Array  
Pandas -> Dask DataFrame  
Scikit-Learn -> Dask-ML  
... -> Dask Futures

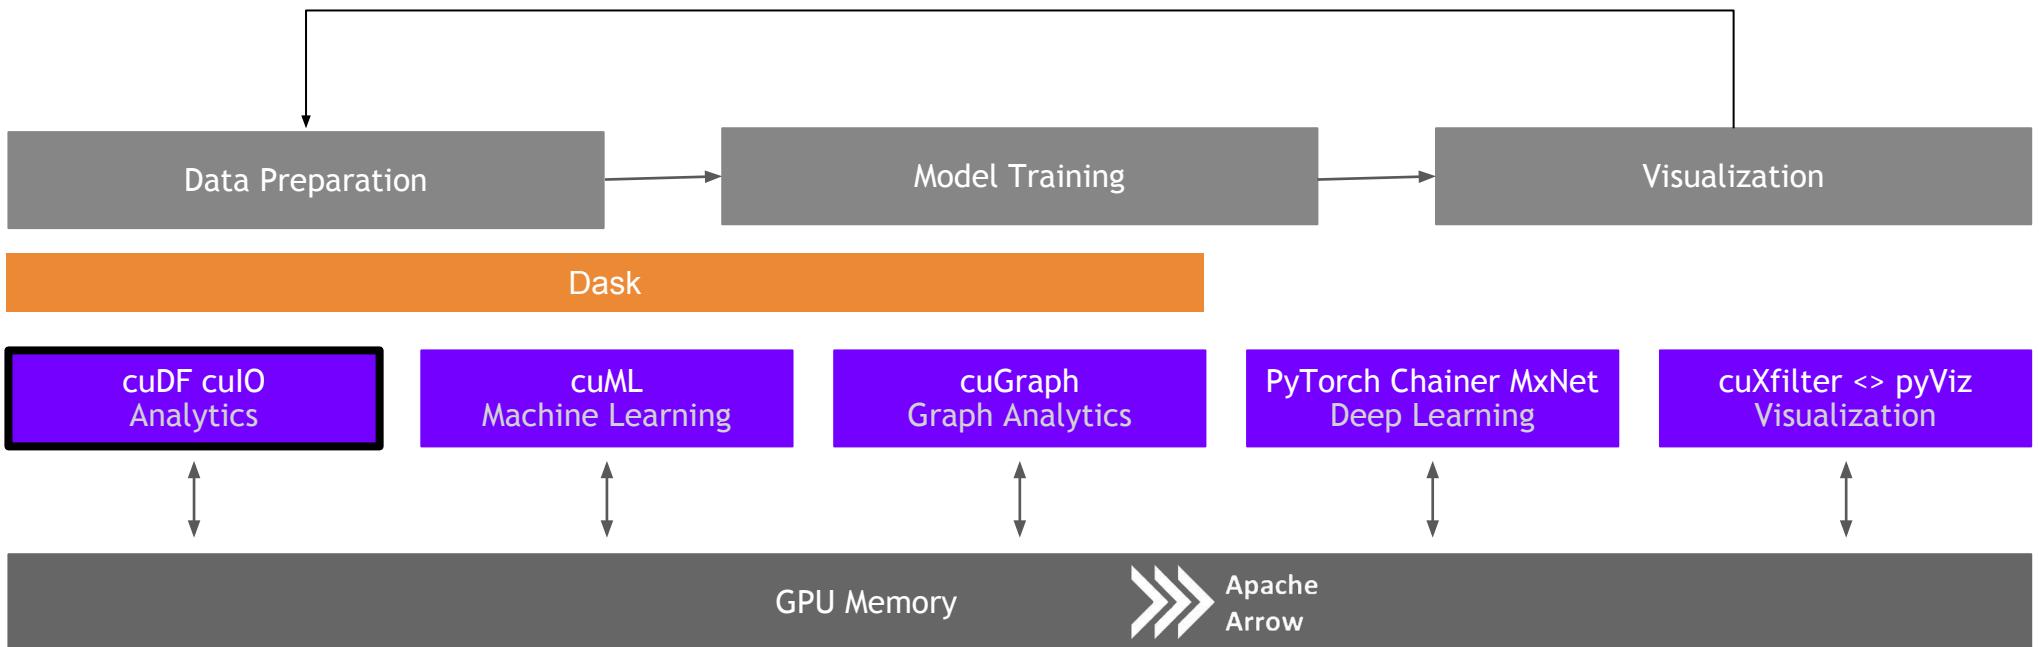


Scale out / Parallelize

cuDF

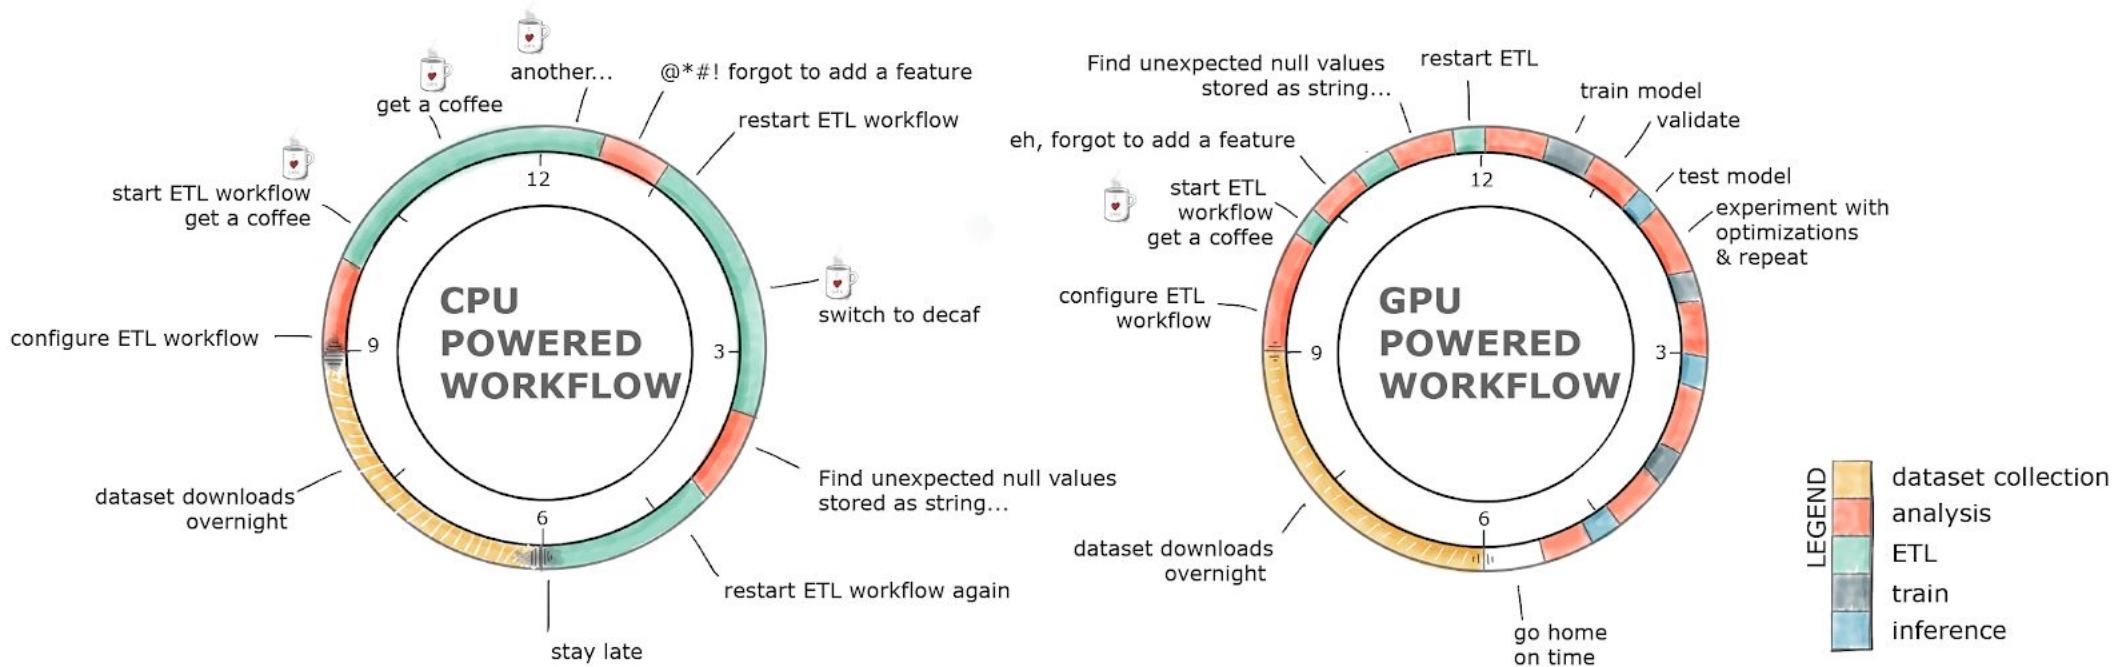
# RAPIDS

## GPU Accelerated data wrangling and feature engineering

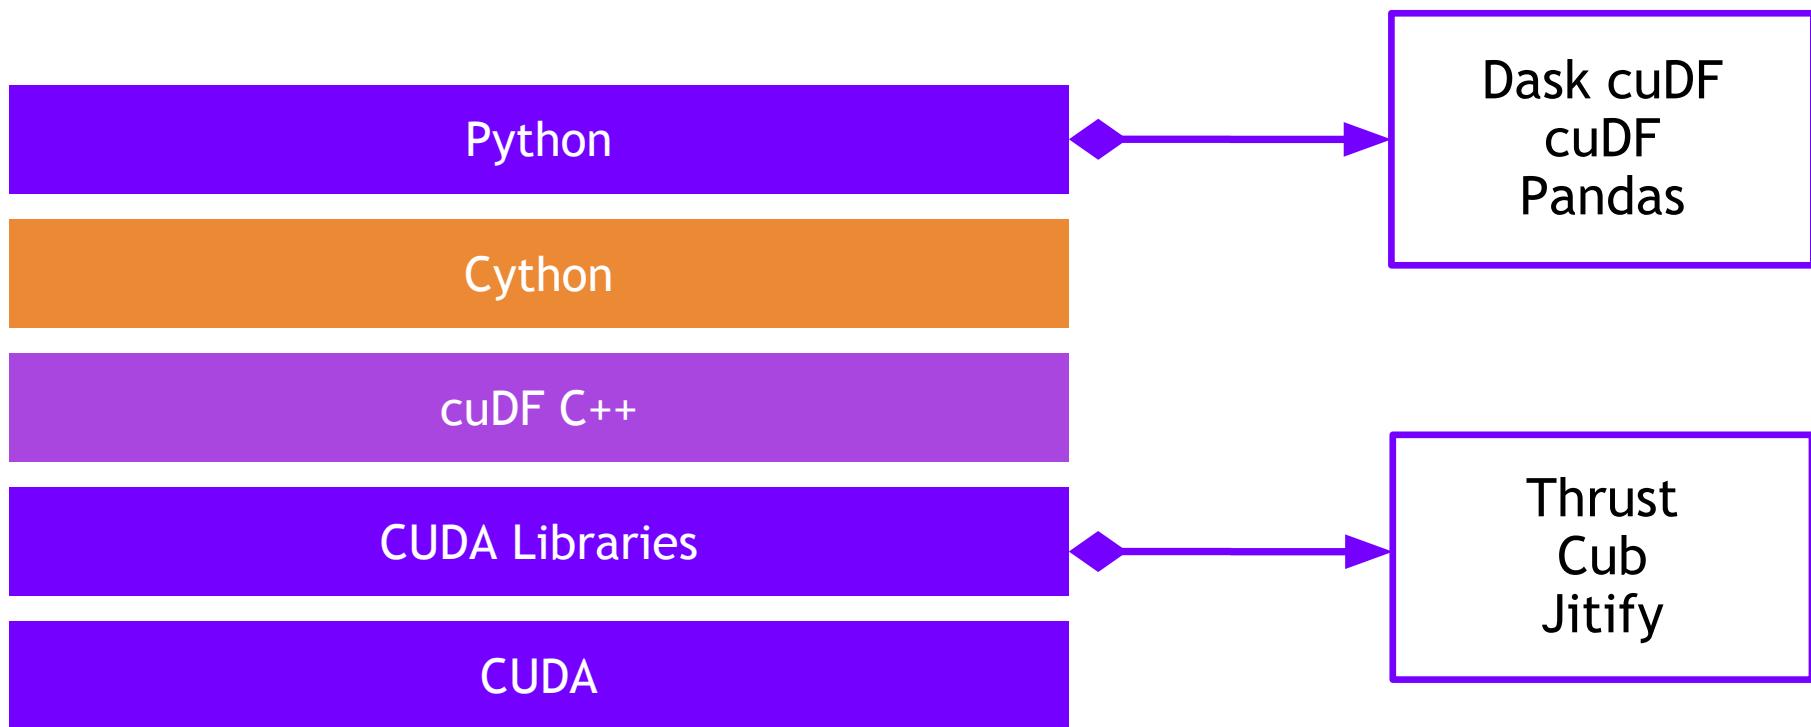


# GPU-Accelerated ETL

The average data scientist spends 90+% of their time in ETL as opposed to training models



# ETL Technology Stack



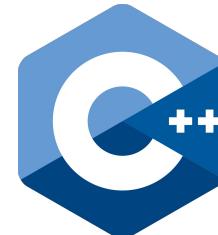
# ETL - the Backbone of Data Science

libcuDF is...

## CUDA C++ Library

- Low level library containing function implementations and C/C++ API
- Importing/exporting Apache Arrow in GPU memory using CUDA IPC
- CUDA kernels to perform element-wise math operations on GPU DataFrame columns
- CUDA sort, join, groupby, reduction, etc. operations on GPU DataFrames

```
void some_function( cudf::column const* input,  
                    cudf::column * output,  
                    args...)  
{  
    // Do something with input  
    // Produce output  
}
```



# ETL - the Backbone of Data Science

## cuDF is...

```
In [2]: #Read in the data. Notice how it decompresses as it reads the data into memory.  
gdf = cudf.read_csv('/rapids/Data/black-friday.zip')
```

```
In [3]: #Taking a look at the data. We use "to_pandas()" to get the pretty printing.  
gdf.head().to_pandas()
```

```
Out[3]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Cat
0	1000001	P00069042	F	0-17	10	A	2	0	3
1	1000001	P00248942	F	0-17	10	A	2	0	1
2	1000001	P00087842	F	0-17	10	A	2	0	12
3	1000001	P00085442	F	0-17	10	A	2	0	12
4	1000002	P00285442	M	55+	16	C	4+	0	8

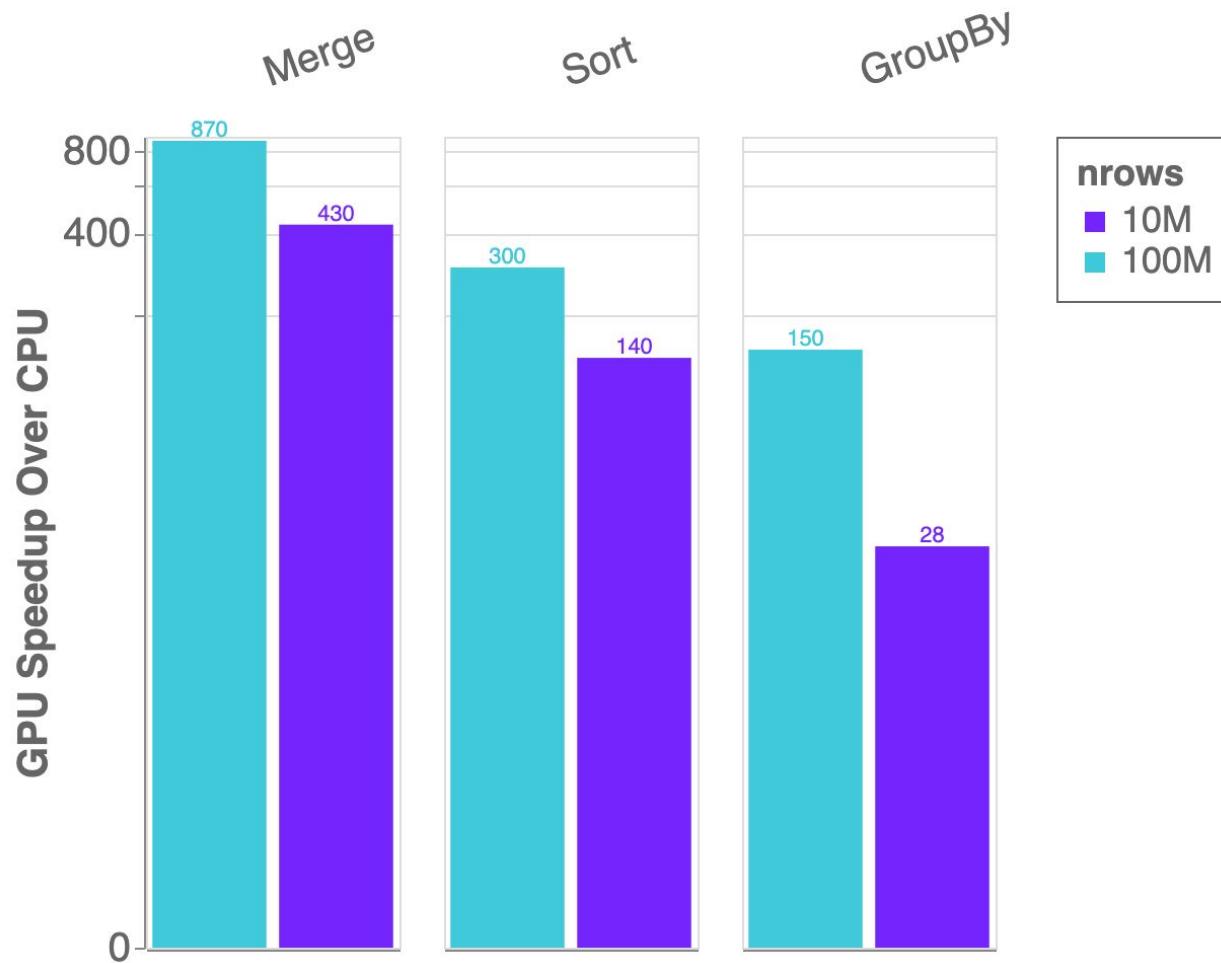
```
In [6]: #grabbing the first character of the years in city string to get rid of plus sign, and converting  
#to int  
gdf['city_years'] = gdf.Stay_In_Current_City_Years.str.get(0).stoi()
```

```
In [7]: #Here we can see how we can control what the value of our dummies with the replace method and turn  
#strings to ints  
gdf['City_Category'] = gdf.City_Category.str.replace('A', '1')  
gdf['City_Category'] = gdf.City_Category.str.replace('B', '2')  
gdf['City_Category'] = gdf.City_Category.str.replace('C', '3')  
gdf['City_Category'] = gdf['City_Category'].str.stoi()
```

## Python Library

- A Python library for manipulating GPU DataFrames following the Pandas API
- Python interface to CUDA C++ library with additional functionality
- Creating GPU DataFrames from Numpy arrays, Pandas DataFrames, and PyArrow Tables
- JIT compilation of User-Defined Functions (UDFs) using Numba

# Benchmarks: single-GPU Speedup vs. Pandas



cuDF v0.10, Pandas 0.24.2

Running on NVIDIA DGX-1:

GPU: NVIDIA Tesla V100 32GB

CPU: Intel(R) Xeon(R) CPU E5-2698 v4  
@ 2.20GHz

Benchmark Setup:

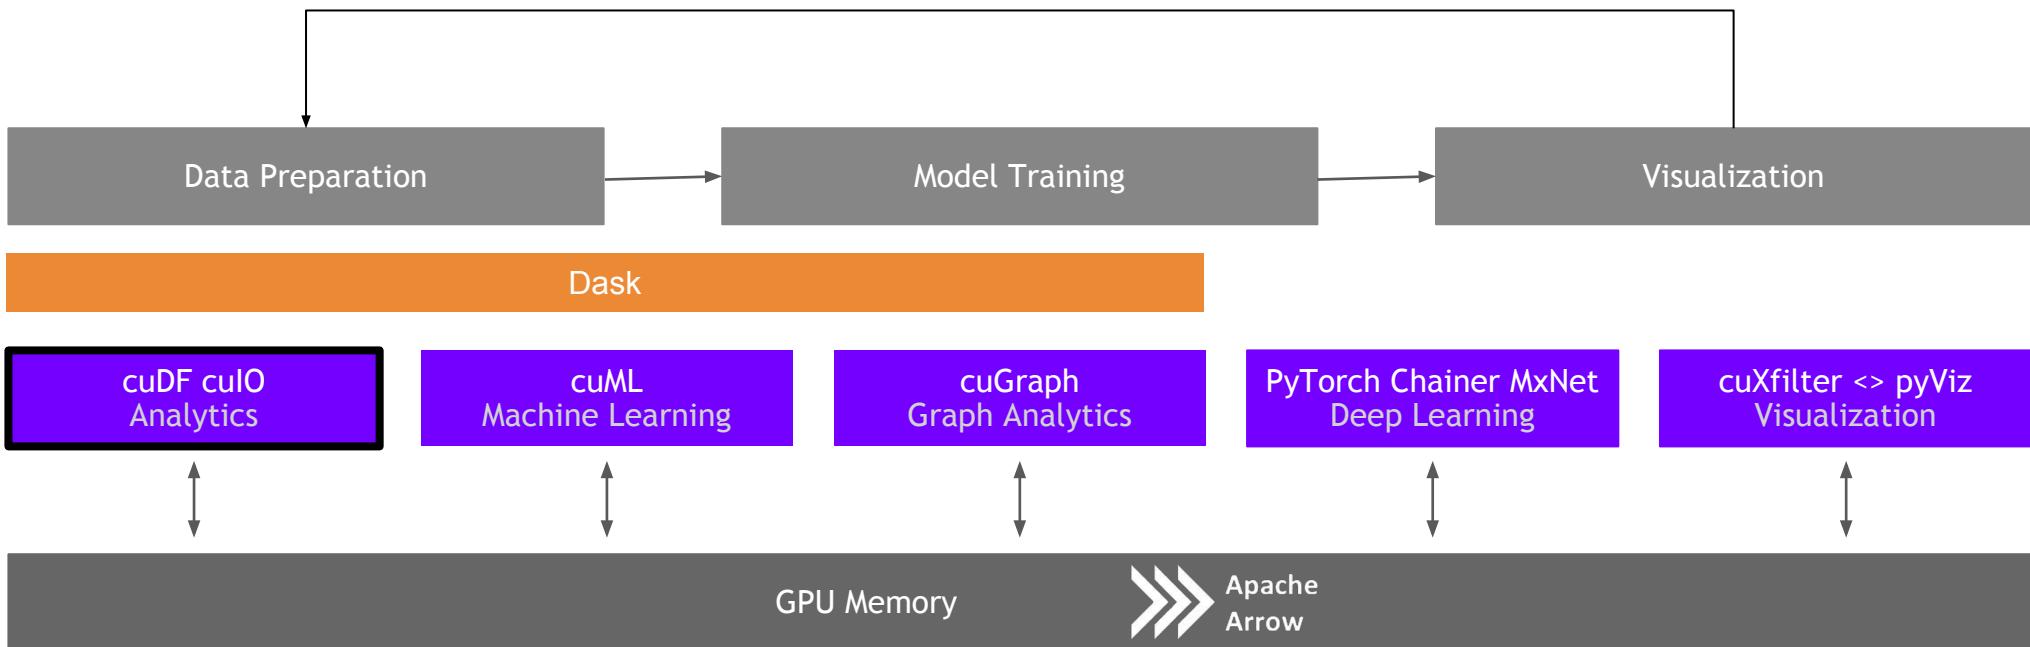
DataFrames: 2x int32 columns key columns,  
3x int32 value columns

Merge: inner

GroupBy: count, sum, min, max calculated  
for each value column

# ETL - the Backbone of Data Science

cuDF is not the end of the story



# ETL - the Backbone of Data Science

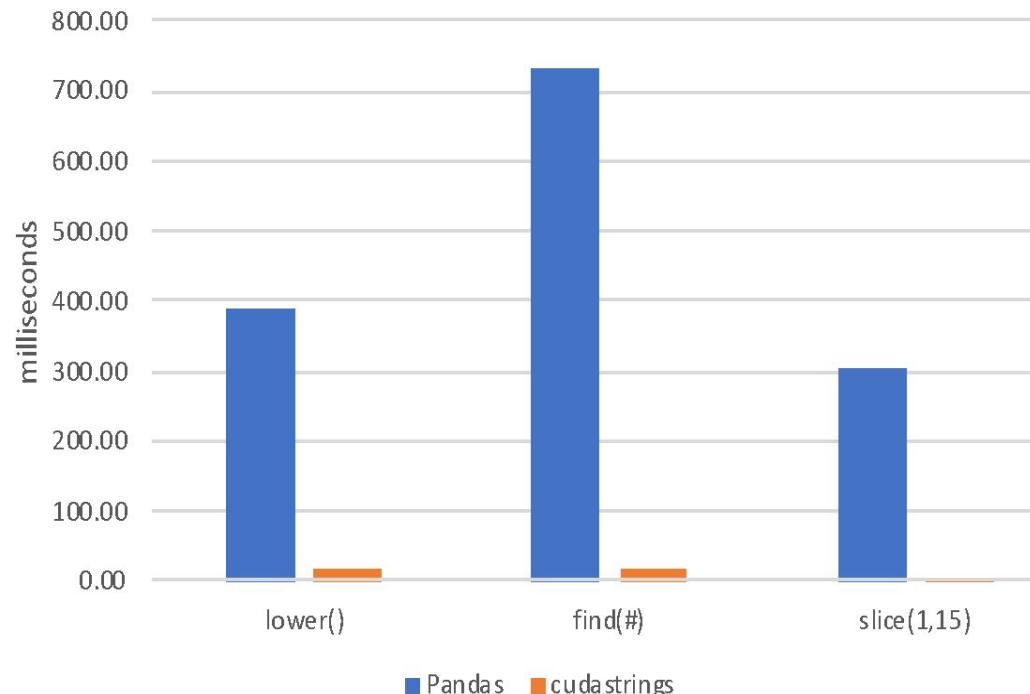
## String Support

### Current v0.10 String Support

- Regular Expressions
- Element-wise operations
  - Split, Find, Extract, Cat, Typecasting, etc...
- String GroupBys, Joins
- Categorical columns fully on GPU
- Combining cuStrings into libcudf

### Future v0.11+ String Support

- Extensive performance optimization
- More Pandas String API compatibility
- JIT-compiled String UDFs
- string columns in libcudf replacing custring



# Extraction is the Cornerstone

## cudf for Faster Data Loading

- Follow Pandas APIs and provide >10x speedup
- CSV Reader - v0.2, CSV Writer v0.8
- Parquet Reader - v0.7, Parquet Writer v0.11
- ORC Reader - v0.7, ORC Writer v0.10
- JSON Reader - v0.8
- Avro Reader - v0.9
- GPU Direct Storage integration in progress for bypassing PCIe bottlenecks!
- Key is GPU-accelerating both parsing and decompression wherever possible

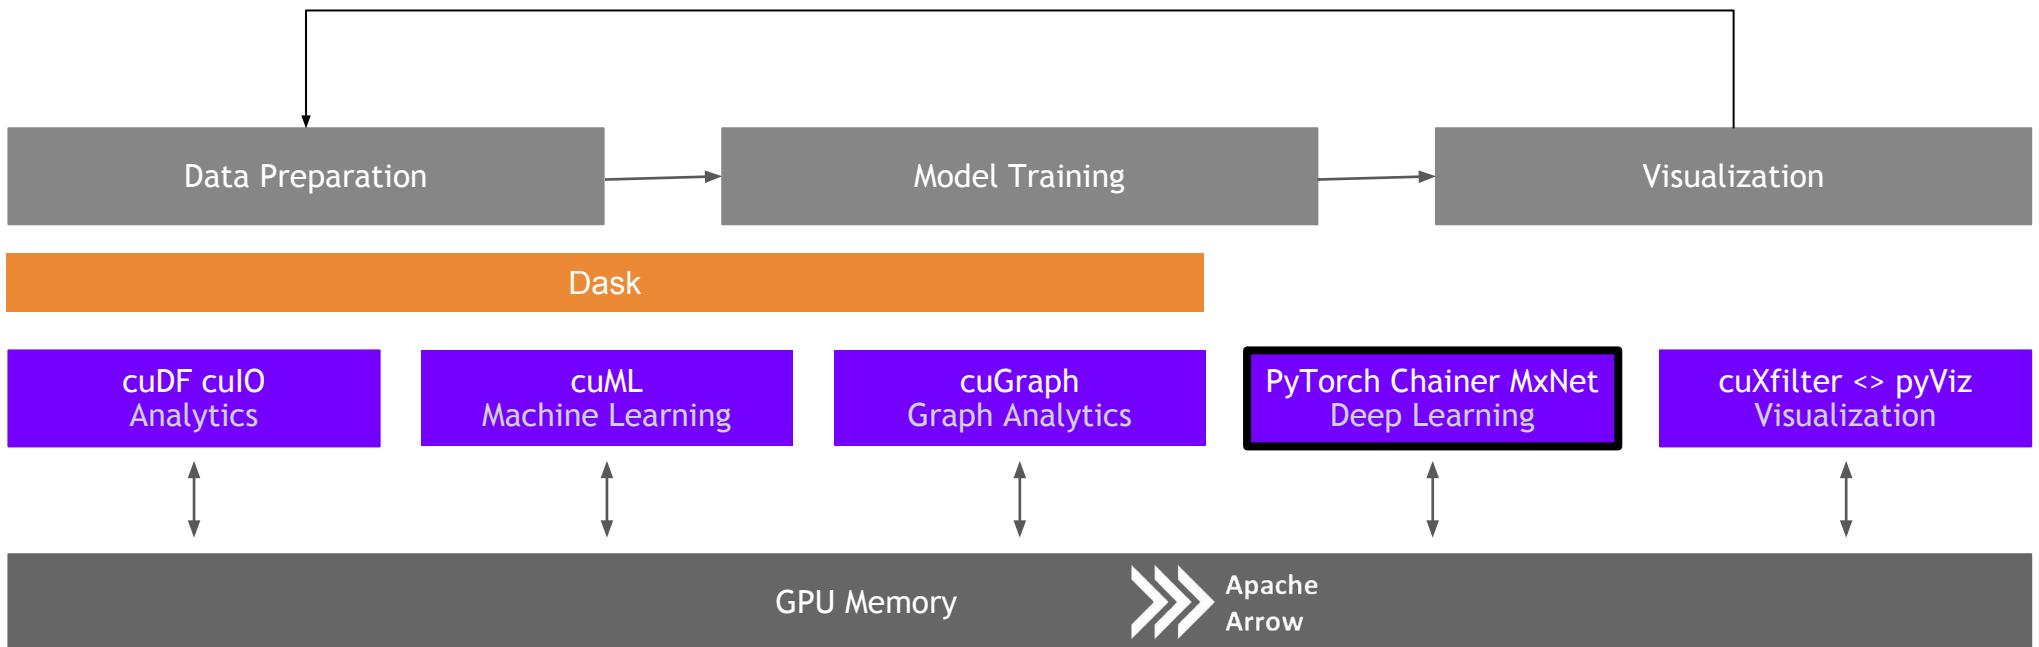
```
1]: import pandas, cudf
2]: %time len(pandas.read_csv('data/nyc/yellow_tripdata_2015-01.csv'))
CPU times: user 25.9 s, sys: 3.26 s, total: 29.2 s
Wall time: 29.2 s
2]: 12748986
3]: %time len(cudf.read_csv('data/nyc/yellow_tripdata_2015-01.csv'))
CPU times: user 1.59 s, sys: 372 ms, total: 1.96 s
Wall time: 2.12 s
3]: 12748986
4]: !du -hs data/nyc/yellow_tripdata_2015-01.csv
1.9G    data/nyc/yellow_tripdata_2015-01.csv
```

Source: Apache Crail blog: [SQL Performance: Part 1 - Input File Formats](#)

ETL is not just DataFrames!

# RAPIDS

## Building bridges into the array ecosystem



# Interoperability for the Win

DLPack and \_\_cuda\_array\_interface\_\_

P Y T  R C H

mpi4py

 m x n e t

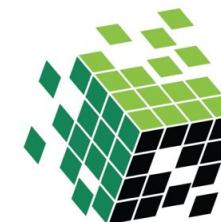




Numba



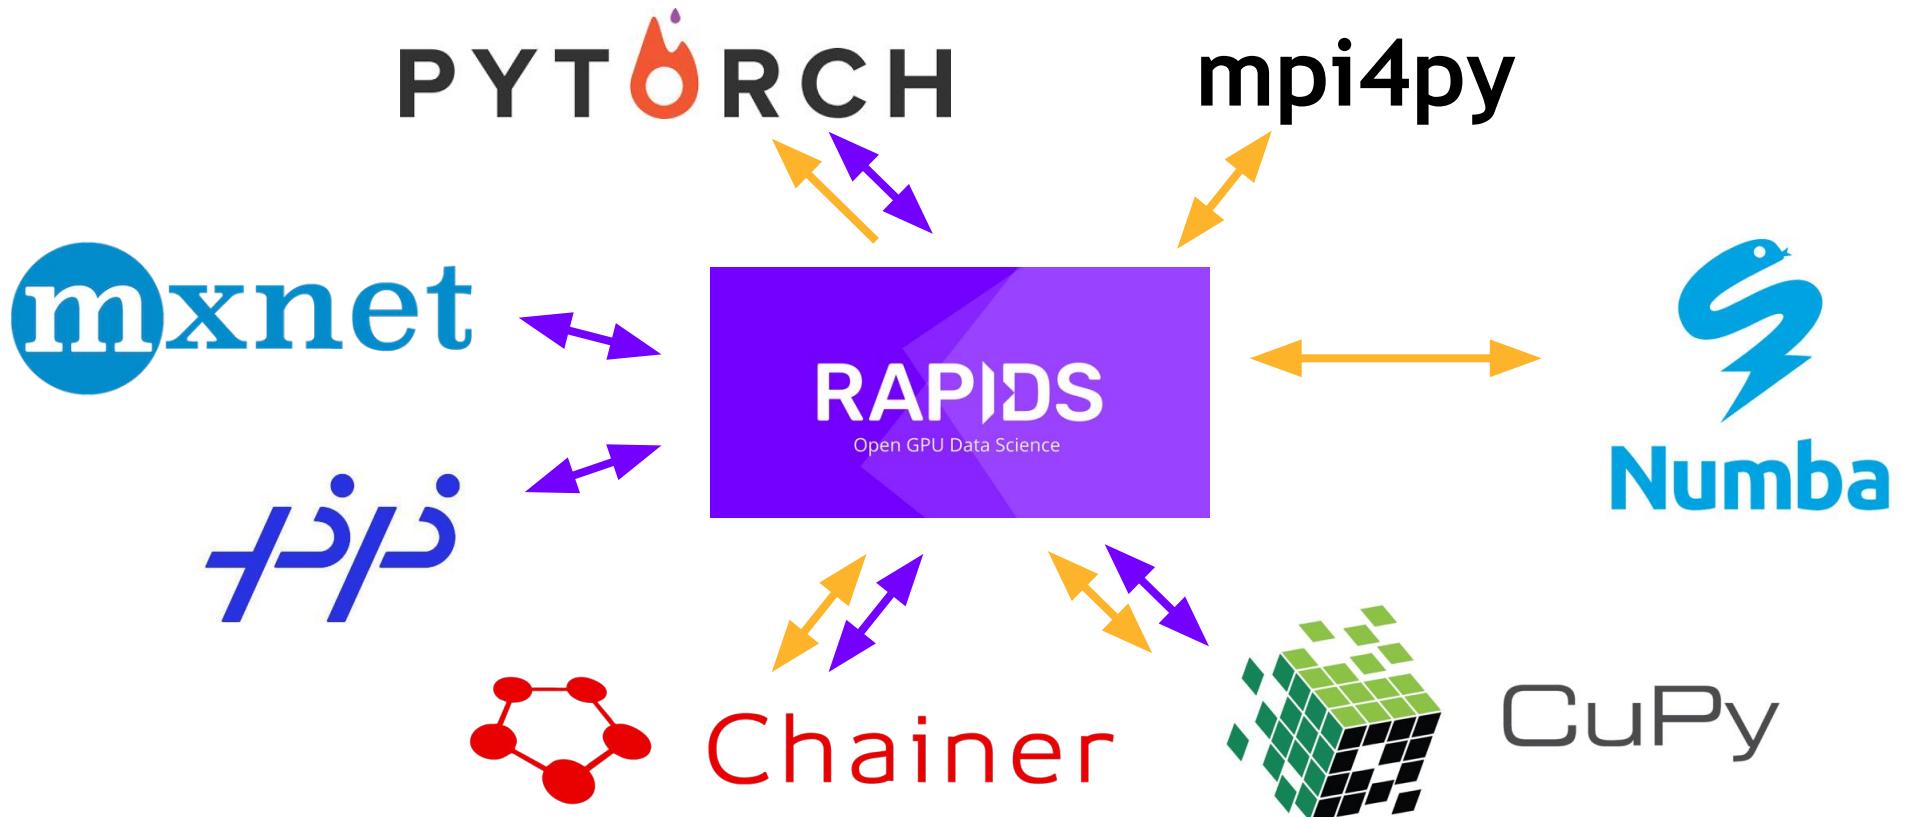
Chainer



CuPy

# Interoperability for the Win

DLPack and `__cuda_array_interface__`



# ETL - Arrays and DataFrames

Dask and CUDA Python arrays



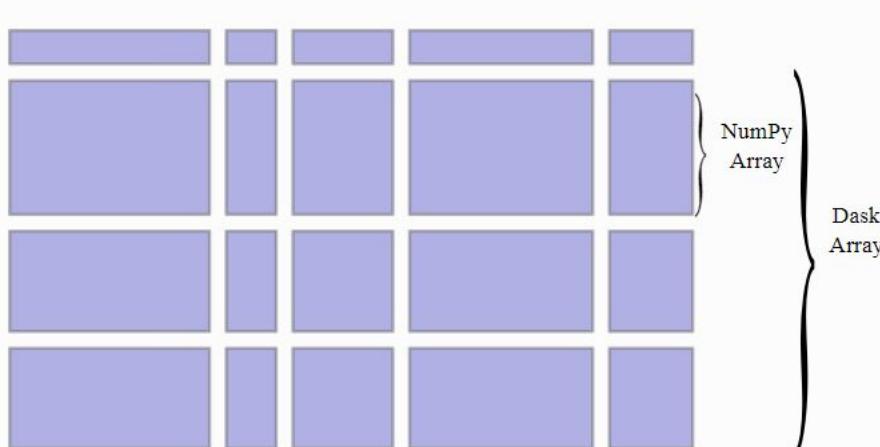
Chainer



CuPy

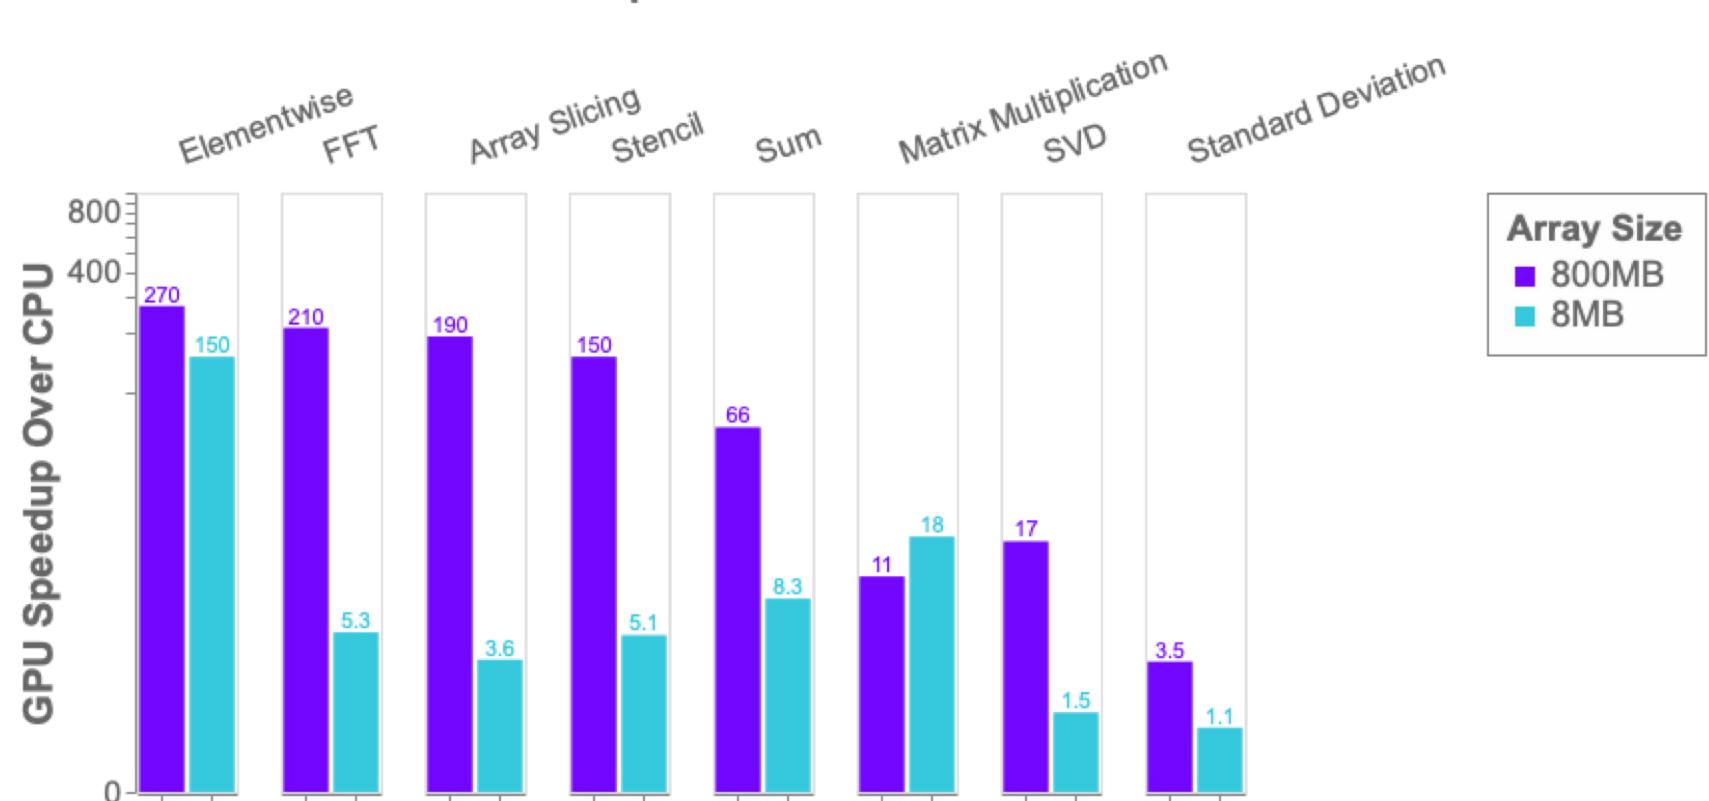


Numba



- Scales NumPy to distributed clusters
- Used in climate science, imaging, HPC analysis up to 100TB size
- Now seamlessly accelerated with GPUs

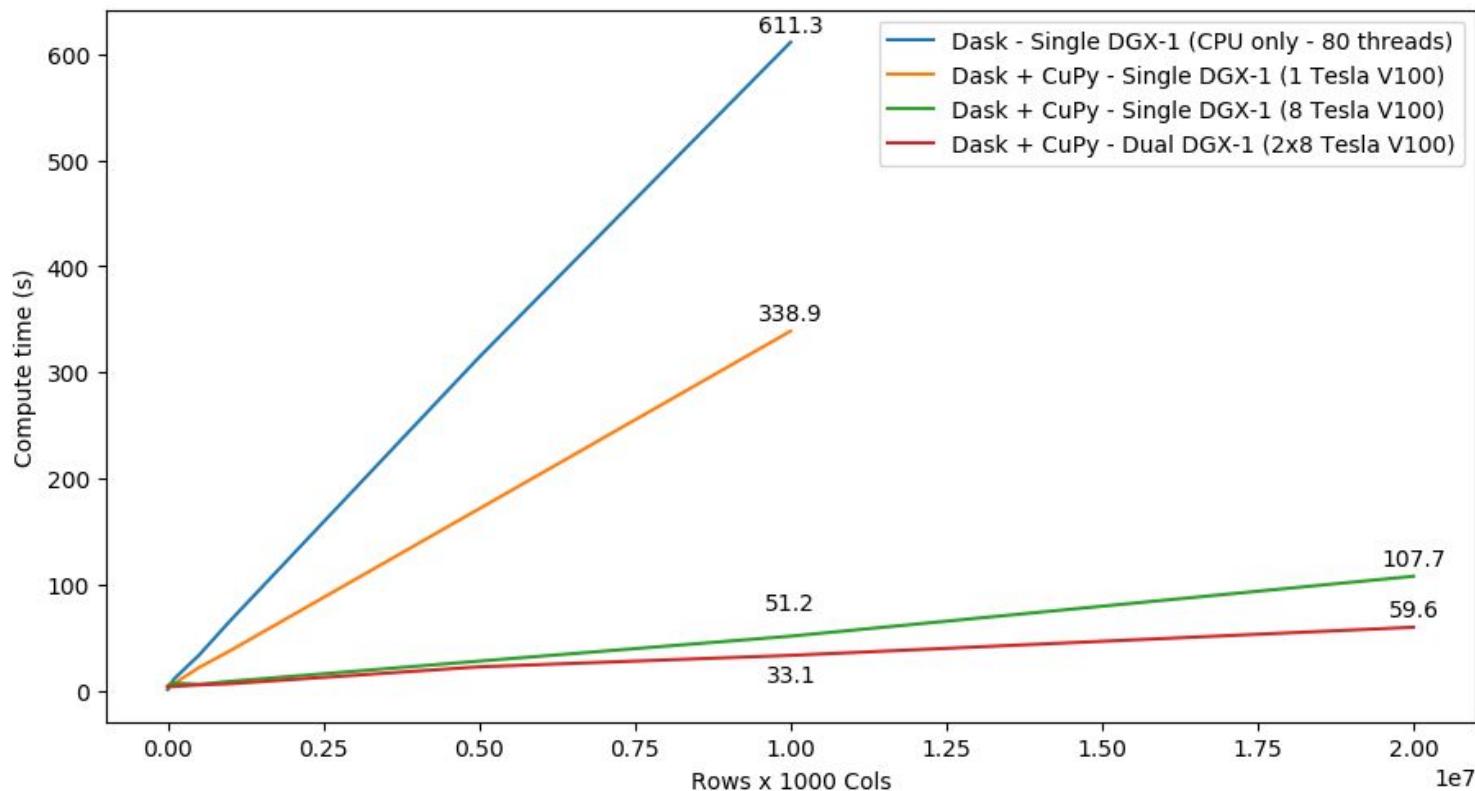
# Benchmark: single-GPU CuPy vs NumPy



More details: <https://blog.dask.org/2019/06/27/single-gpu-cupy-benchmarks>

# SVD Benchmark

## Dask and CuPy Doing Complex Workflows



# Also...Achievement Unlocked:

## Petabyte Scale Data Analytics with Dask and CuPy

Architecture	Time
Single CPU Core	2hr 39min
Forty CPU Cores	11min 30s
One GPU	1min 37s
Eight GPUs	19s

<https://blog.dask.org/2019/01/03/dask-array-gpus-first-steps>



### 3.2 PETABYTES IN LESS THAN 1 HOUR

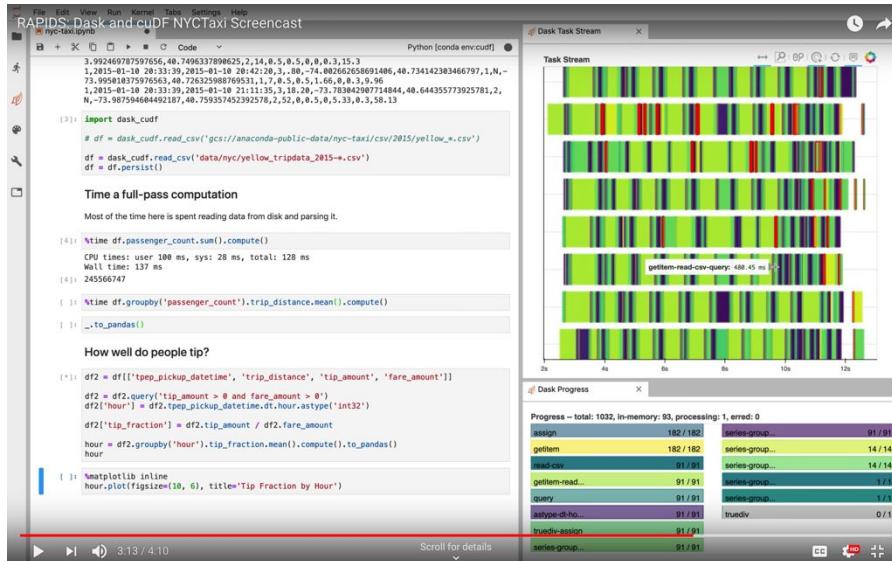
Distributed GPU array | parallel reduction | using 76x GPUs

Array size	Wall Time (data creation + compute)
3.2 PB (20M x 20M doubles)	54 min 51 s

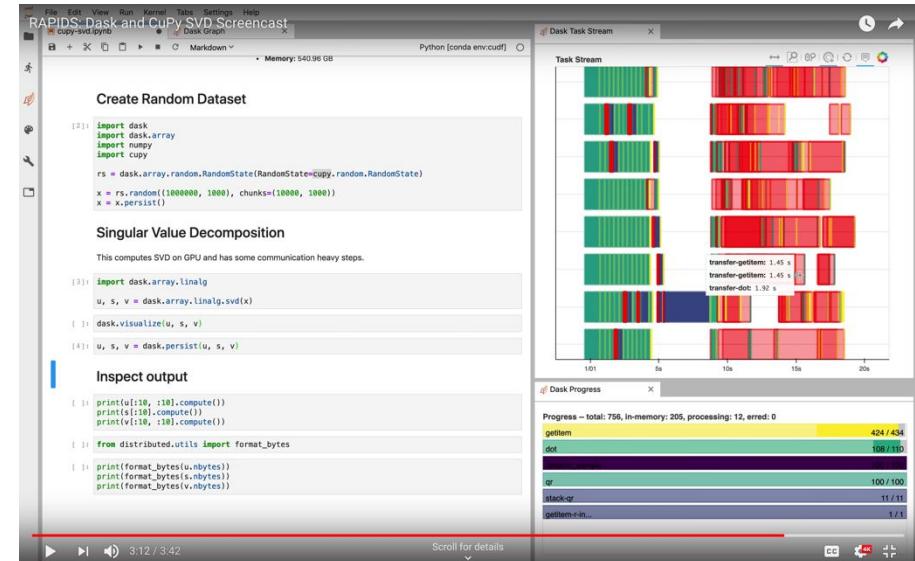
**Cluster configuration:** 20x GCP instances, each instance has:  
**CPU:** 1 VM socket (Intel Xeon CPU @ 2.30GHz), 2-core, 2 threads/core, 132GB mem, GbE ethernet, 950 GB disk  
**GPU:** 4x NVIDIA Tesla P100-16GB-PCIe (total GPU DRAM across nodes 1.22 TB)  
**Software:** Ubuntu 18.04, RAPIDS 0.5.1, Dask=1.1.1, Dask-Distributed=1.1.1, CuPY=5.2.0, CUDA 10.0.130

# ETL - Arrays and DataFrames

## More Dask Awesomeness from RAPIDS



<https://youtu.be/gV0cykgsTPM>

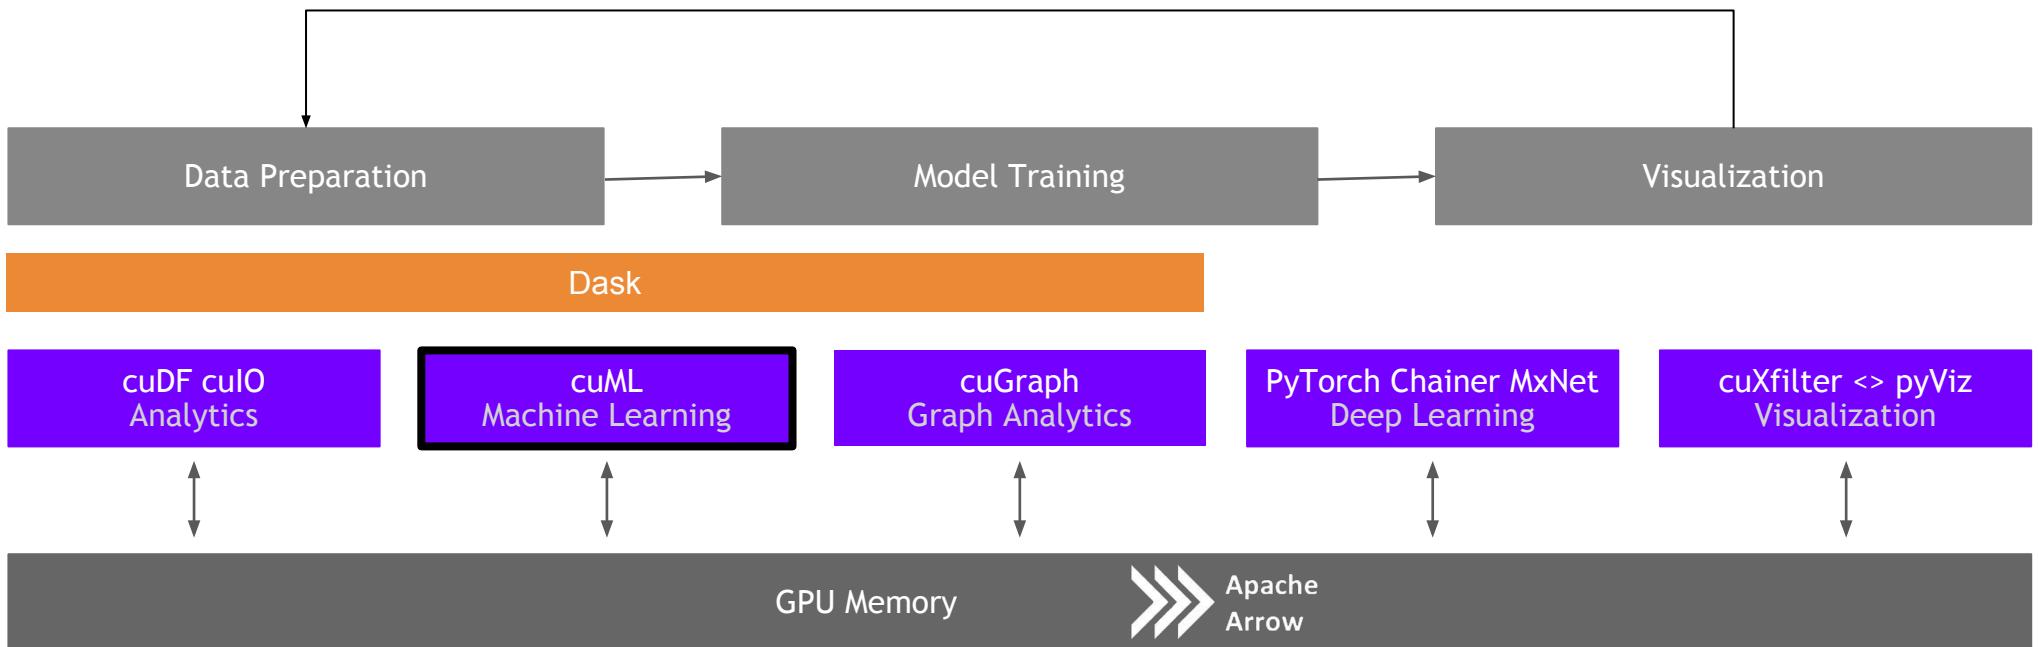


[https://youtu.be/R5CiXti\\_MWo](https://youtu.be/R5CiXti_MWo)

cuML

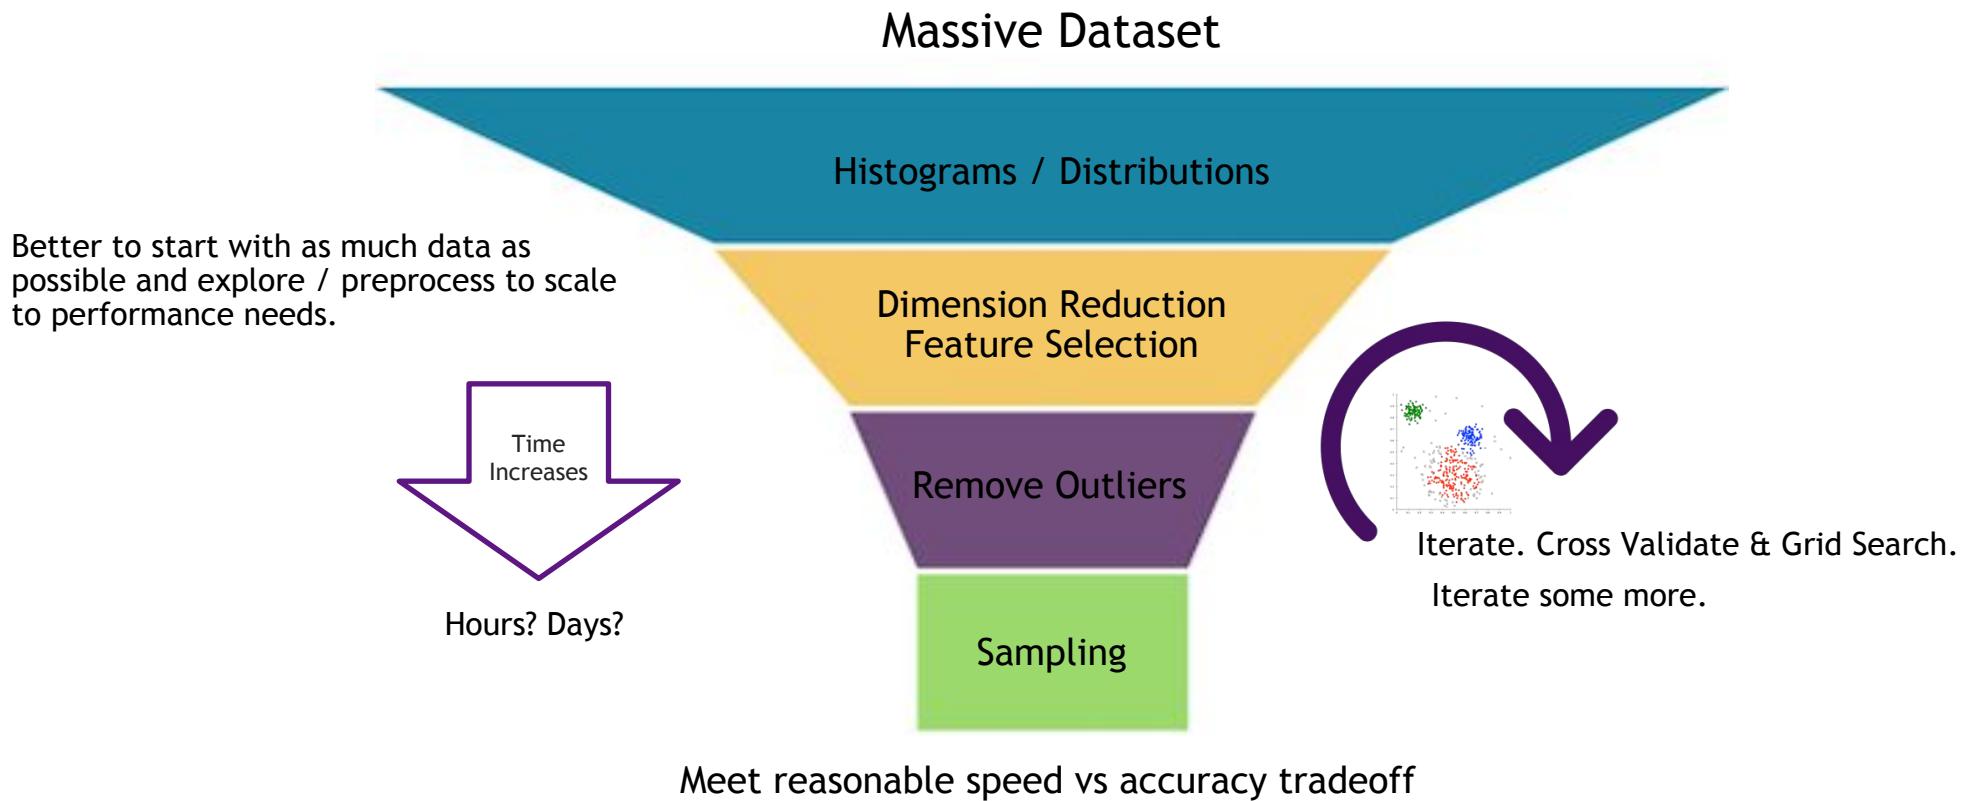
# Machine Learning

## More models more problems

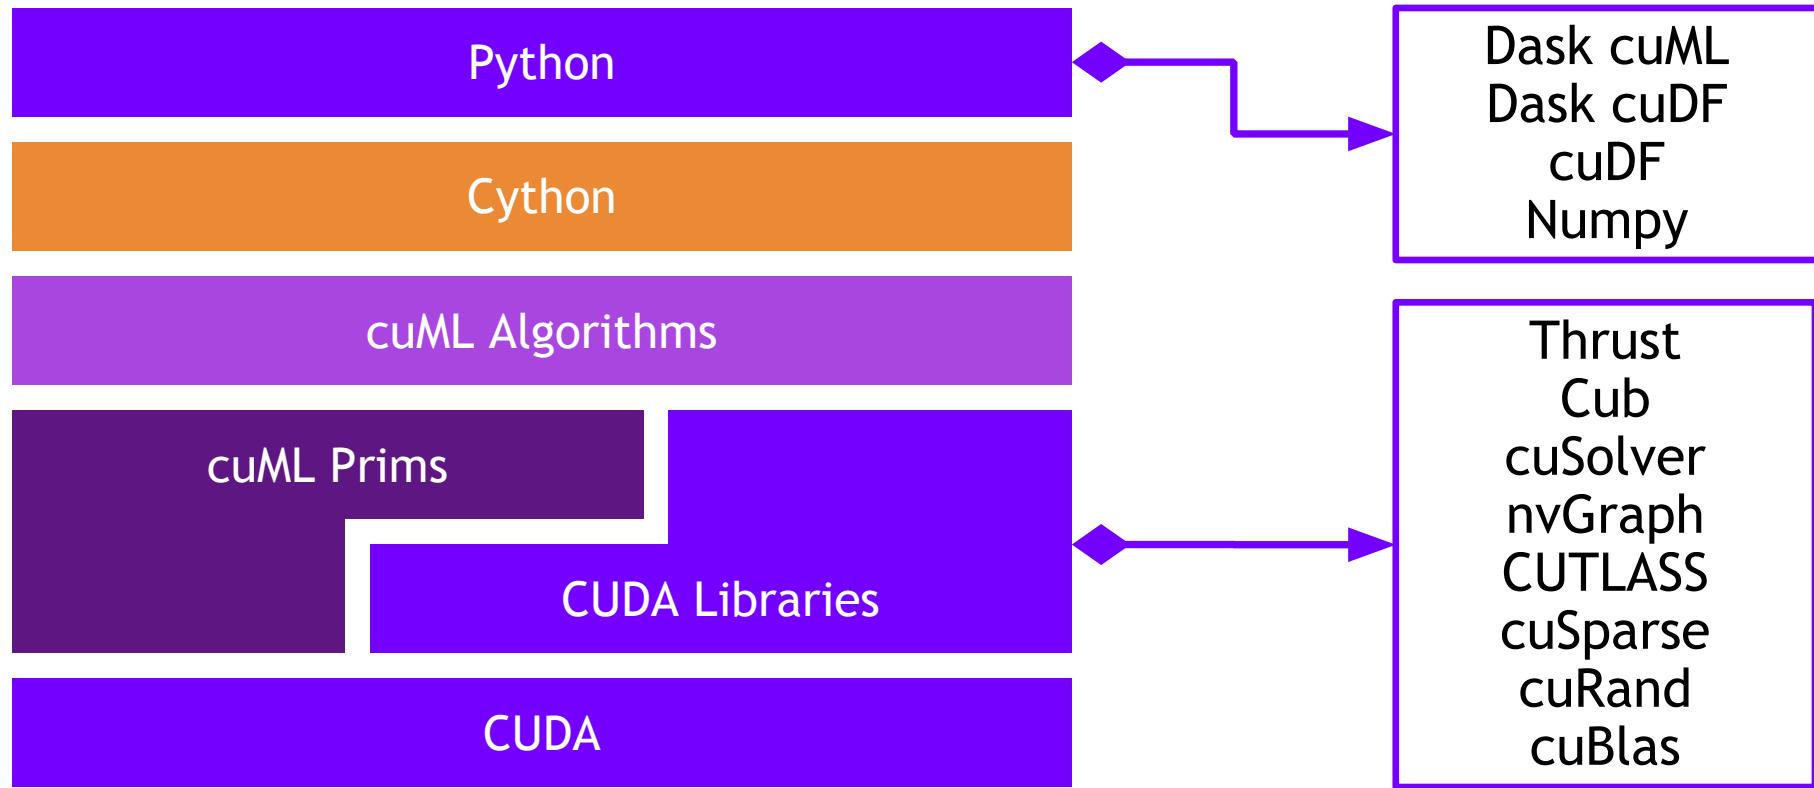


# Problem

Data sizes continue to grow

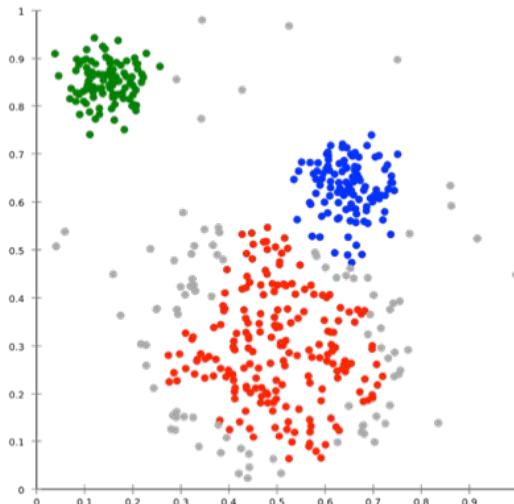


# ML Technology Stack



# Algorithms

## GPU-accelerated Scikit-Learn



Cross Validation

Hyper-parameter Tuning

More to come!

Classification / Regression

Inference

Clustering

Decomposition & Dimensionality Reduction

Time Series

Decision Trees / Random Forests

Linear Regression

Logistic Regression

K-Nearest Neighbors

**Support Vector Machine Classification**

Random forest / GBDT inference

K-Means  
DBSCAN  
Spectral Clustering

Principal Components  
Singular Value Decomposition  
UMAP  
Spectral Embedding  
**T-SNE**

Holt-Winters  
Kalman Filtering

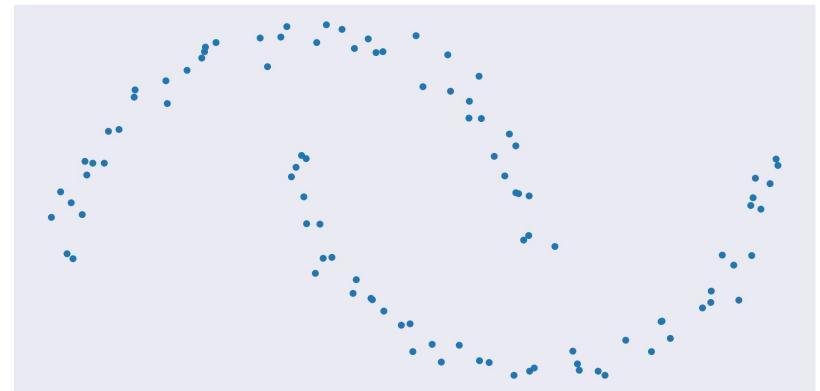
Key:

- Preexisting
- **NEW for 0.10**

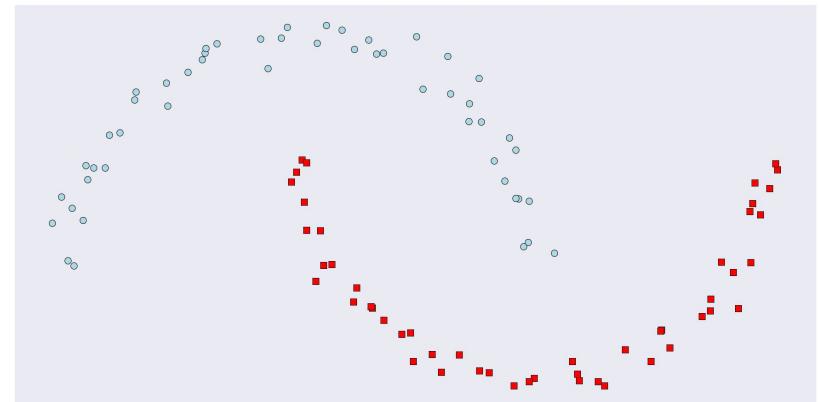
# RAPIDS matches common Python APIs

## CPU-Based Clustering

```
from sklearn.datasets import make_moons  
import pandas  
  
X, y = make_moons(n_samples=int(1e2),  
                   noise=0.05, random_state=0)  
  
X = pandas.DataFrame({'fea%d'%i: X[:, i]  
                      for i in range(X.shape[1])})
```



```
from sklearn.cluster import DBSCAN  
dbSCAN = DBSCAN(eps = 0.3, min_samples = 5)  
  
dbSCAN.fit(X)  
  
y_hat = dbSCAN.predict(X)
```



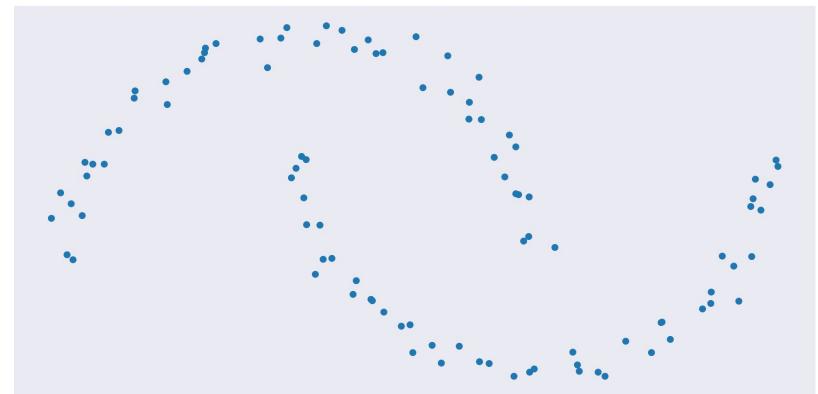
# RAPIDS matches common Python APIs

## GPU-Accelerated Clustering

```
from sklearn.datasets import make_moons
import cudf

X, y = make_moons(n_samples=int(1e2),
                   noise=0.05, random_state=0)

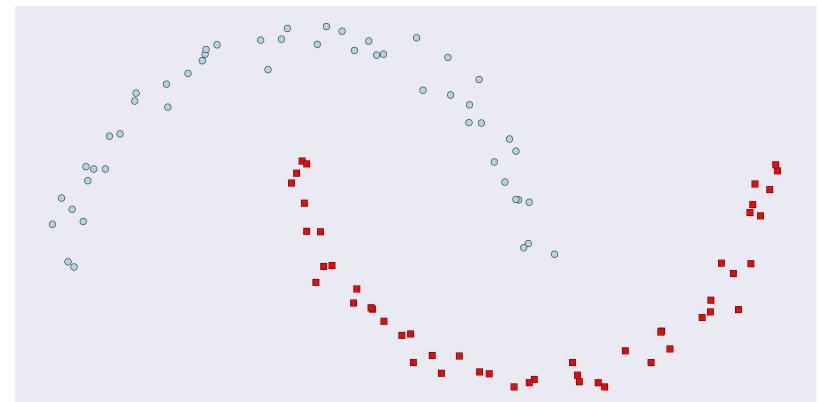
X = cudf.DataFrame({'fea%d'%i: X[:, i]
                    for i in range(X.shape[1])})
```



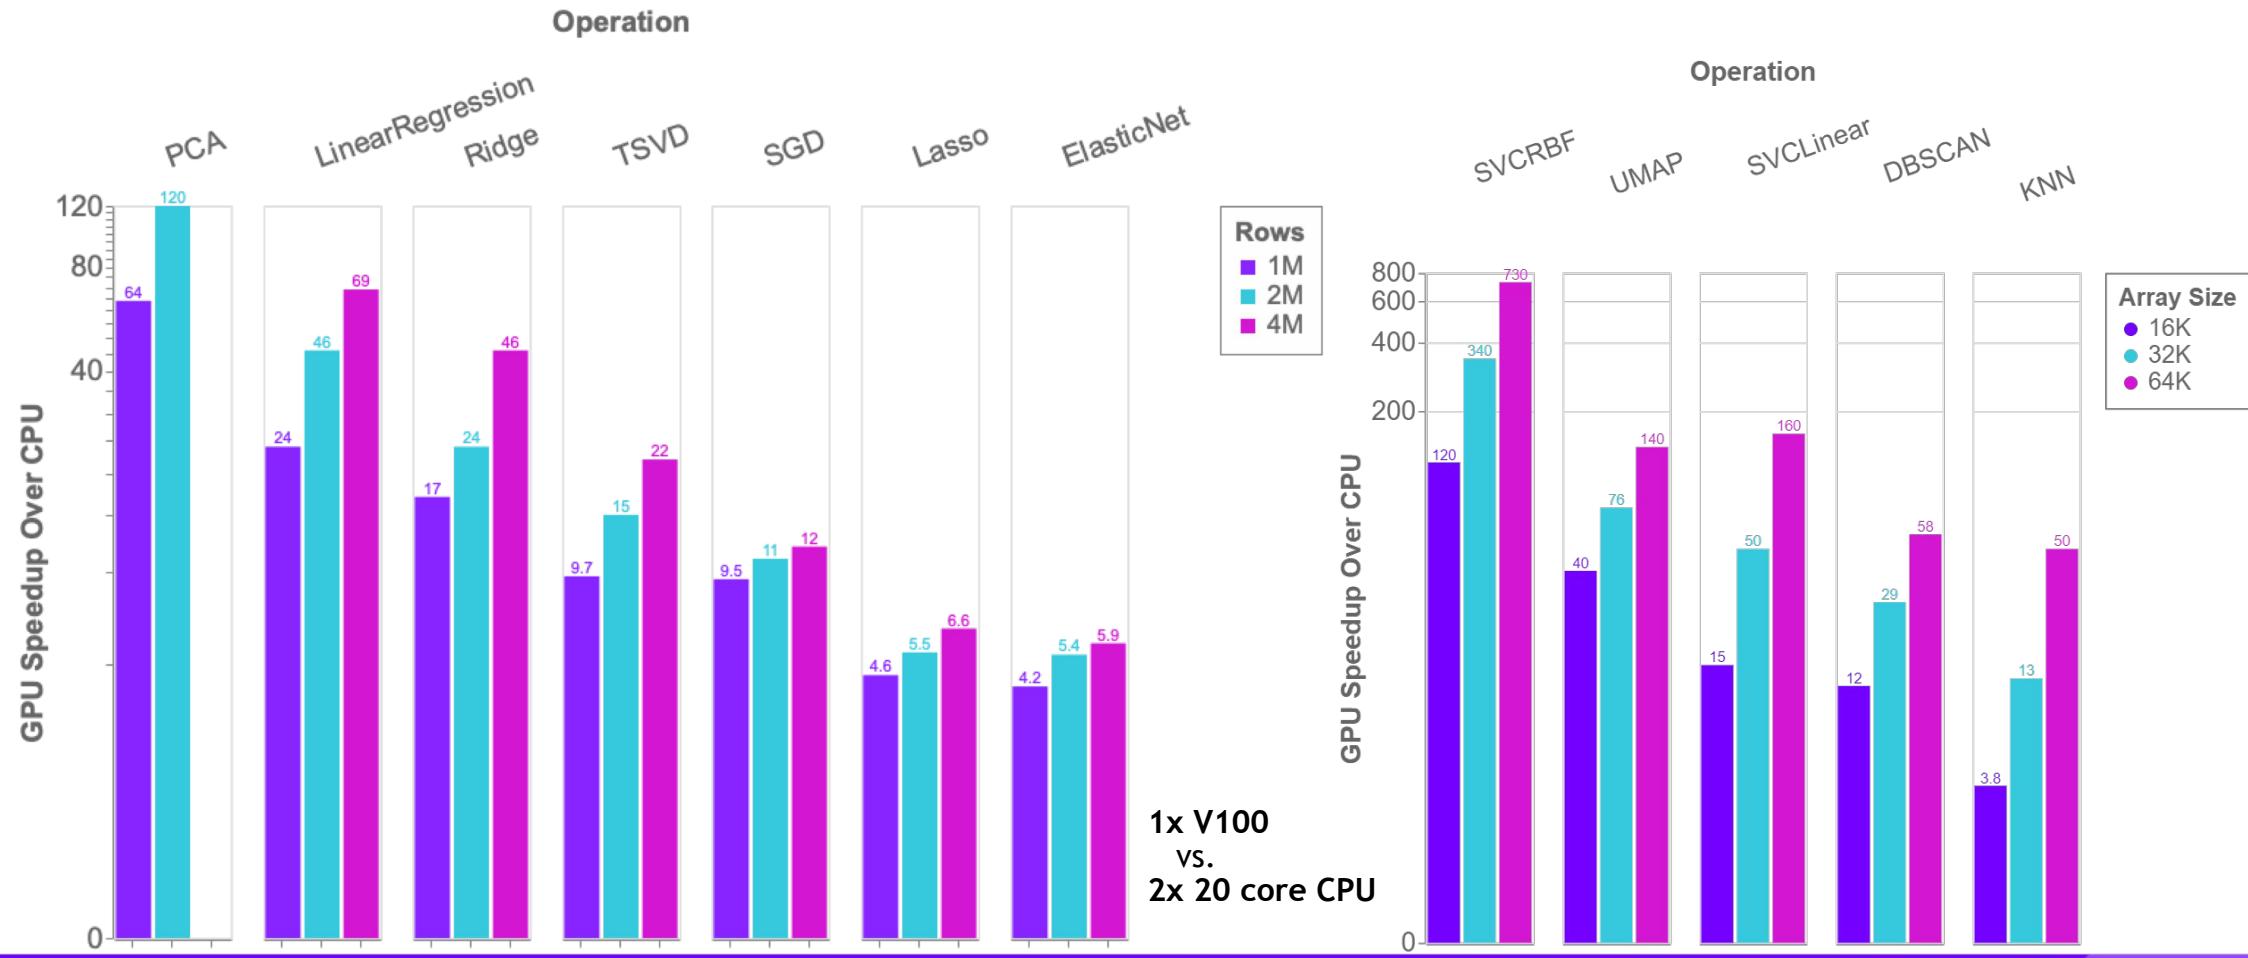
```
from cuml import DBSCAN
dbscan = DBSCAN(eps = 0.3, min_samples = 5)

dbscan.fit(X)

y_hat = dbscan.predict(X)
```



# Benchmarks: single-GPU cuML vs scikit-learn

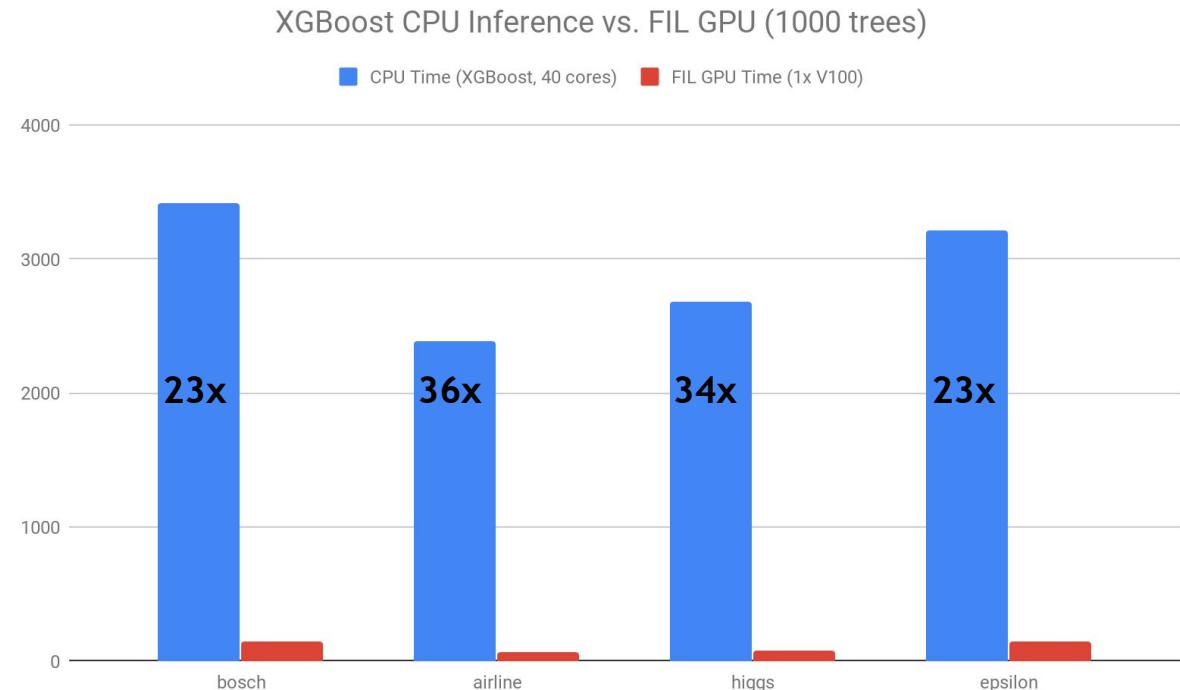


# Forest Inference

Taking models from training to production

cuML's **Forest Inference Library** accelerates prediction (inference) for random forests and boosted decision trees:

- Works with existing saved models (XGBoost and LightGBM today, scikit-learn RF and cuML RF soon)
- Lightweight Python API
- Single V100 GPU can infer up to 34x faster than XGBoost dual-CPU node
- Over 100 million forest inferences per sec (with 1000 trees) on a DGX-1



# Road to 1.0

## October 2019 - RAPIDS 0.10

cuML	Single-GPU	Multi-GPU	Multi-Node-Multi-GPU
Gradient Boosted Decision Trees (GBDT)			
GLM			
Logistic Regression			
Random Forest			
K-Means			
K-NN			
DBSCAN			
UMAP			
Holt-Winters			
Kalman Filter			
t-SNE			
Principal Components			
Singular Value Decomposition			
SVM			

# Road to 1.0

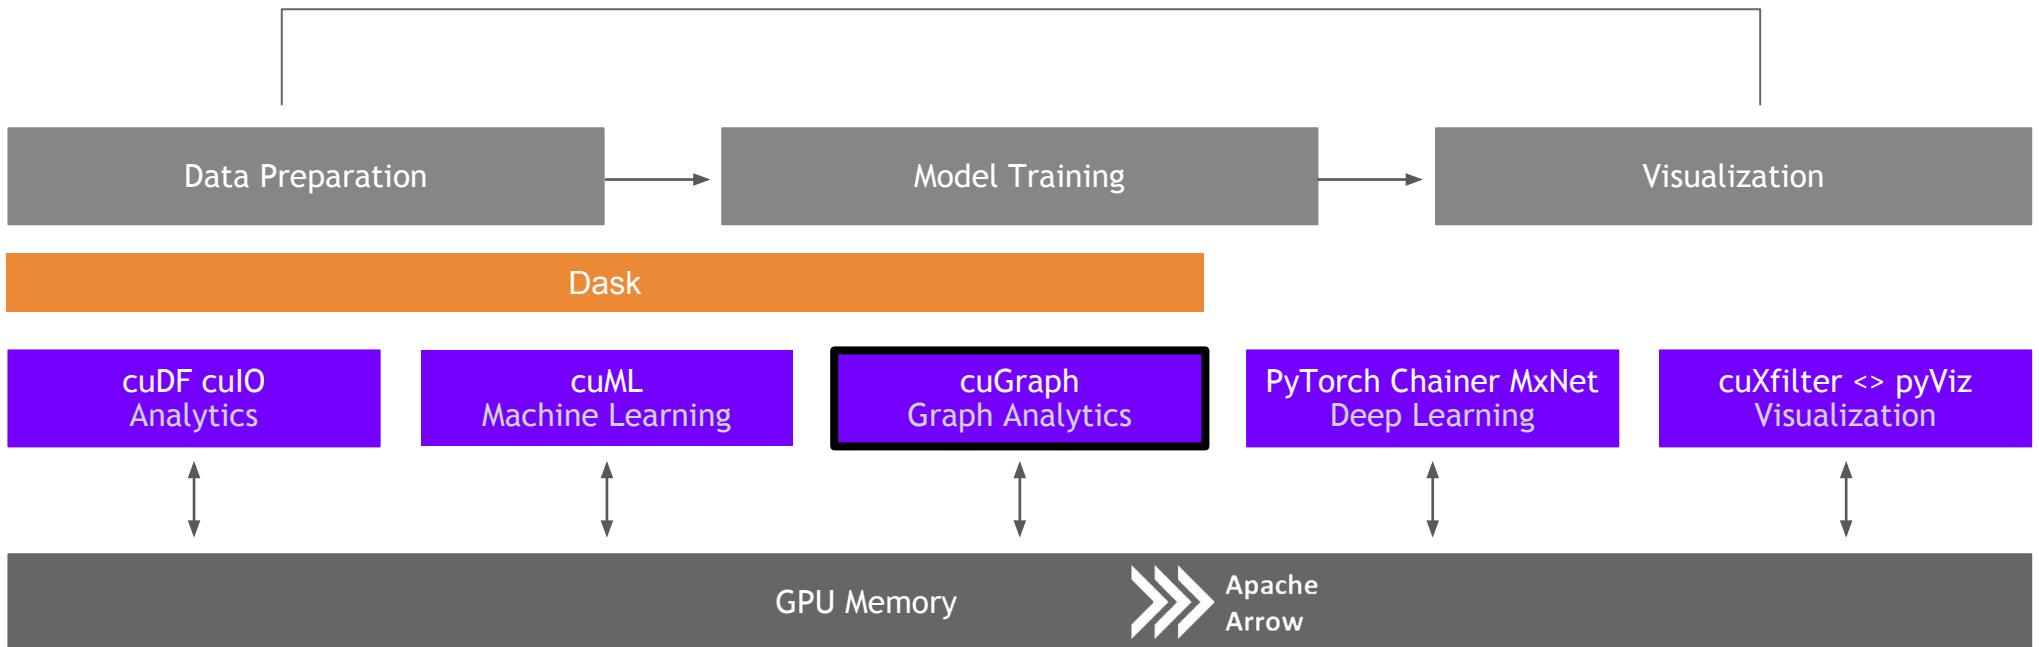
## March 2020 - RAPIDS 0.13

cuML	Single-GPU	Multi-Node-Multi-GPU
Gradient Boosted Decision Trees (GBDT)		
GLM		
Logistic Regression		
Random Forest		
K-Means		
K-NN		
DBSCAN		
UMAP		
ARIMA & Holt-Winters		
Kalman Filter		
t-SNE		
Principal Components		
Singular Value Decomposition		
SVM		

# cuGraph

# Graph Analytics

## More connections more insights



# GOALS AND BENEFITS OF CUGRAPH

Focus on Features and User Experience

## Breakthrough Performance

- Up to 500 million edges on a single 32GB GPU
- Multi-GPU support for scaling into the billions of edges

## Multiple APIs

- Python: Familiar NetworkX-like API
- C/C++: lower-level granular control for application developers

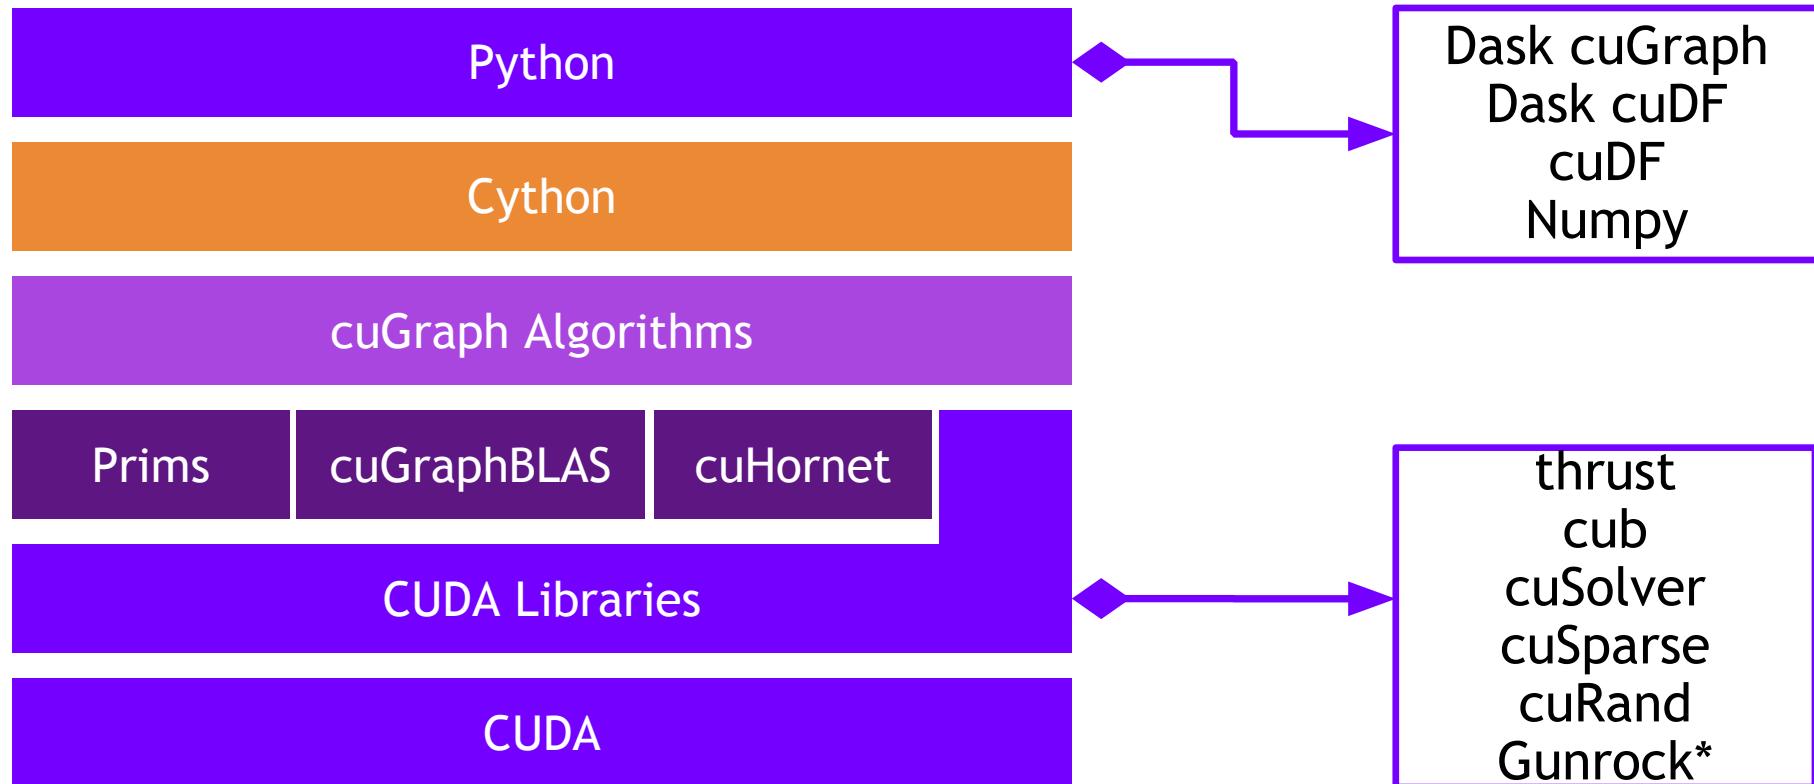
## Seamless Integration with cuDF and cuML

- Property Graph support via DataFrames

## Growing Functionality

- Extensive collection of algorithm, primitive, and utility functions

# Graph Technology Stack

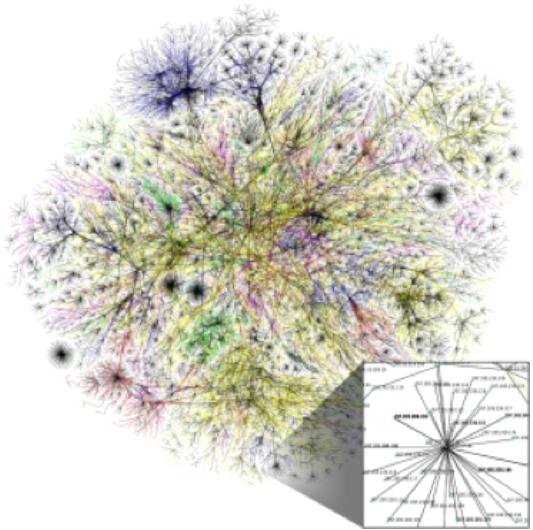


nvGRAPH has been Opened Sourced and integrated into cuGraph. A legacy version is available in a RAPIDS GitHub repo

\* Gunrock is from UC Davis

# Algorithms

## GPU-accelerated NetworkX



Renumbering

More to come!

Query Language

Multi-GPU

Utilities

Community

Components

Link Analysis

Link Prediction

Traversal

Structure

Spectral Clustering  
Balanced-Cut  
Modularity Maximization  
Louvain  
Subgraph Extraction  
**KCore**

Weakly Connected Components  
Strongly Connected Components

Page Rank (Multi-GPU)  
Personal Page Rank  
Katz

Jaccard  
Weighted Jaccard  
Overlap Coefficient

Single Source Shortest Path (SSSP)  
Breadth First Search (BFS)

Triangle Counting  
COO-to-CSR (Multi-GPU)  
Transpose

# Louvain Single Run

```
G = cugraph.Graph()  
G.add_edge_list(gdf["src_0"], gdf["dst_0"], gdf["data"])  
df, mod = cugraph.nvLouvain(G)
```

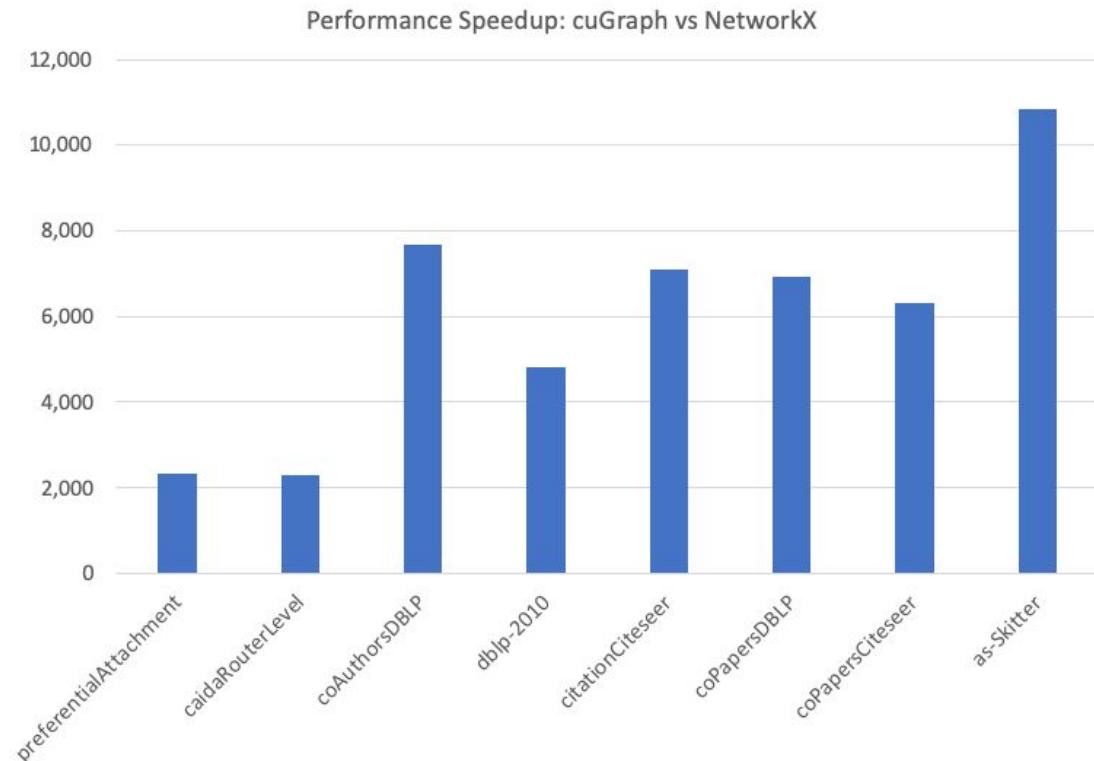
**Louvain returns:**

cudf.DataFrame with two names columns:

louvain["vertex"]: The vertex id.

louvain["partition"]: The assigned partition.

Dataset	Nodes	Edges
preferentialAttachment	100,000	999,970
caidaRouterLevel	192,244	1,218,132
coAuthorsDBLP	299,067	299,067
dblp-2010	326,186	1,615,400
citationCiteseer	268,495	2,313,294
coPapersDBLP	540,486	30,491,458
coPapersCiteseer	434,102	32,073,440
as-Skitter	1,696,415	22,190,596



# Multi-GPU PageRank Performance

PageRank portion of the HiBench benchmark suite

HiBench Scale	Vertices	Edges	CSV File (GB)	# of GPUs	# of CPU Threads	PageRank for 3 Iterations (secs)
Huge	5,000,000	198,000,000	3	1		1.1
BigData	50,000,000	1,980,000,000	34	3		5.1
BigData x2	100,000,000	4,000,000,000	69	6		9.0
BigData x4	200,000,000	8,000,000,000	146	12		18.2
BigData x8	400,000,000	16,000,000,000	300	16		31.8
BigData x8	400,000,000	16,000,000,000	300		800*	5760*

\*BigData x8, 100x 8-vCPU nodes, Apache Spark GraphX ⇒ 96 mins!

# Road to 1.0

## October 2019 - RAPIDS 0.10

cuGraph	Single-GPU	Multi-GPU	Multi-Node-Multi-GPU
Jaccard and Weighted Jaccard			
Page Rank			
Personal Page Rank			
SSSP			
BFS			
Triangle Counting			
Subgraph Extraction			
Katz Centrality			
Betweenness Centrality			
Connected Components (Weak and Strong)			
Louvain			
Spectral Clustering			
K-Cores			

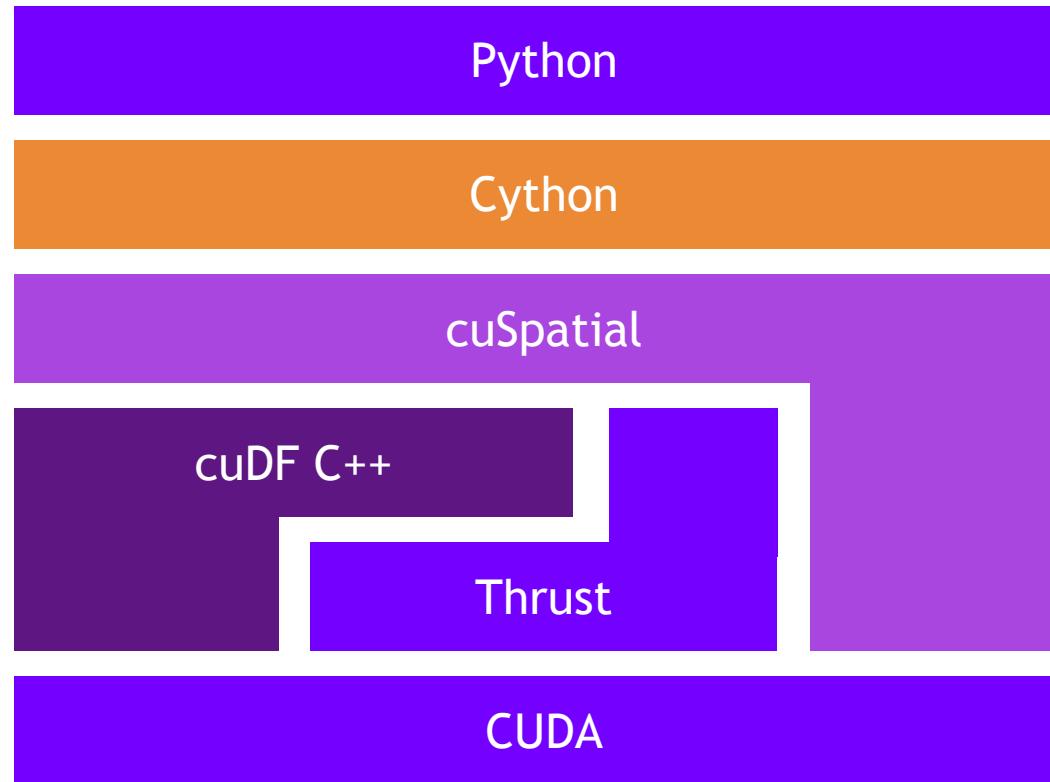
# Road to 1.0

## March 2020 - RAPIDS 0.13

cuGraph	Single-GPU	Multi-Node-Multi-GPU
Jaccard and Weighted Jaccard		
Page Rank		
Personal Page Rank		
SSSP		
BFS		
Triangle Counting		
Subgraph Extraction		
Katz Centrality		
Betweenness Centrality		
Connected Components (Weak and Strong)		
Louvain		
Spectral Clustering		
K-Cores		

cuSpatial

# cuSpatial Technology Stack



# cuSpatial 0.10

## Breakthrough Performance & Ease of Use

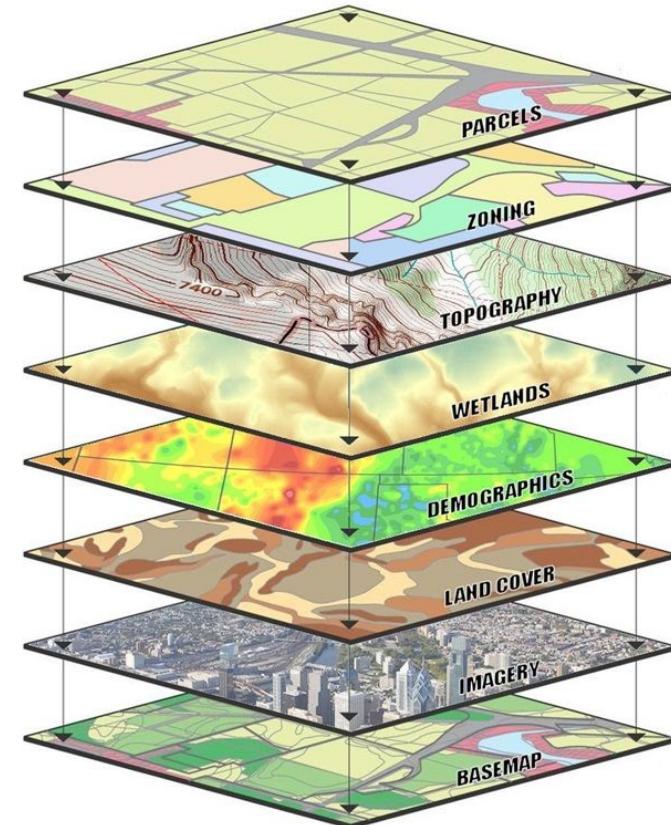
- Up to 1000x faster than CPU spatial libraries
- Python and C++ APIs for maximum usability and integration

## Growing Functionality

- Extensive collection of algorithm, primitive, and utility functions for spatial analytics

## Seamless Integration into RAPIDS

- cuDF for data loading, cuGraph for routing optimization, and cuML for clustering are just a few examples



# cuSpatial

## 0.10 and Beyond

Layer	0.10/0.11 Functionality	Functionality Roadmap (2020)
High-level Analytics	C++ Library w. Python bindings enabling distance, speed, trajectory similarity, trajectory clustering	C++ Library w. Python bindings for additional spatio-temporal trajectory clustering, acceleration, dwell-time, salient locations, trajectory anomaly detection, origin destination, etc.
Graph layer	cuGraph	Map matching, Djikstra algorithm, Routing
Query layer	Spatial Window	Nearest Neighbor, KNN, Spatiotemporal range search and joins
Index layer		Grid, Quad Tree, R-Tree, Geohash, Voronoi Tessellation
Geo-operations	Point in polygon (PIP), Haversine distance, Hausdorff distance, lat-lon to xy transformation	Line intersecting polygon, Other distance functions, Polygon intersection, union
Geo-representation	Shape primitives, points, polylines, polygons	Additional shape primitives

# cuSpatial 0.10

## Performance at a Glance

cuSpatial Operation	Input data	cuSpatial Runtime	Reference Runtime	Speedup
Point-in-Polygon Test	1.3+ million vehicle point locations and 27 Region of Interests	1.11 ms (C++) 1.50 ms (Python) [Nvidia Titan V]	334 ms (C++, optimized serial) 130468.2 ms (python Shapely API, serial) [Intel i7-7800X]	301X (C++) 86,978X (Python)
Haversine Distance Computation	13+ million Monthly NYC taxi trip pickup and drop-off locations	7.61 ms (Python) [Nvidia T4]	416.9 ms (Numba) [Nvidia T4]	54.7X (Python)
Hausdorff Distance Computation (for clustering)	10,700 trajectories with 1.3+ million points	13.5s [Quadro V100]	19227.5s (Python SciPy API, serial) [Intel i7-6700K]	1,400X (Python)

# Community

# Ecosystem Partners

## CONTRIBUTORS



## ADOPTERS

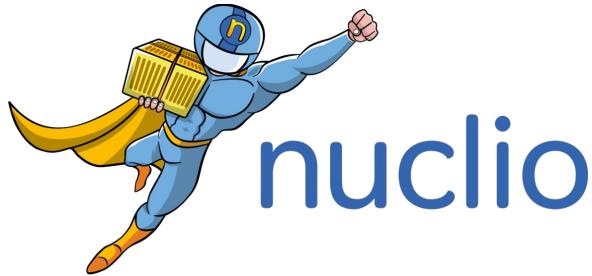


## OPEN SOURCE



# Building on top of RAPIDS

A bigger, better, stronger ecosystem for all



High-Performance  
Serverless event and  
data processing that  
utilizes RAPIDS for GPU  
Acceleration

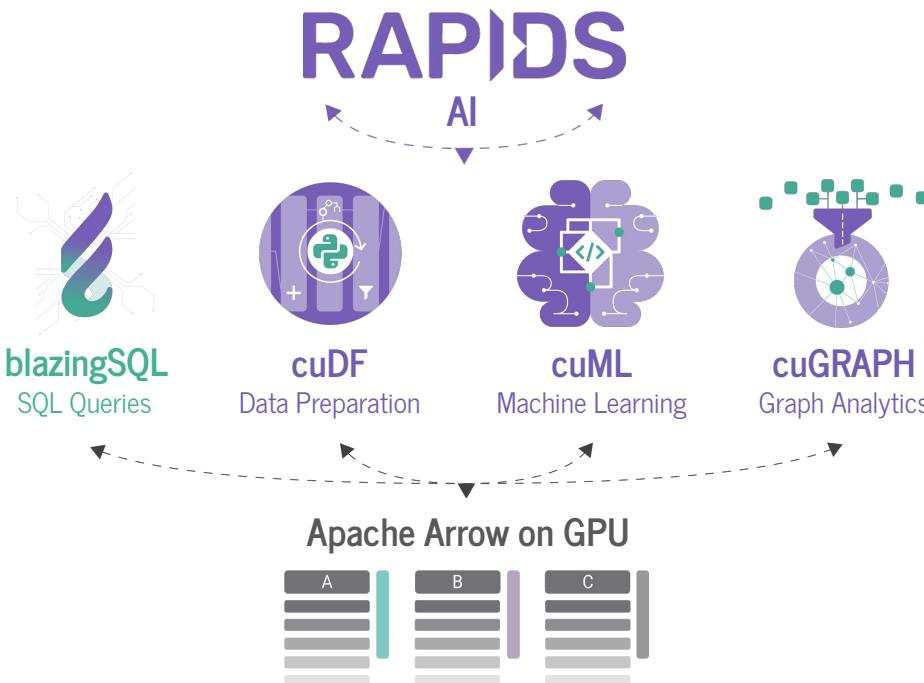


GPU accelerated SQL  
engine built on top of  
RAPIDS

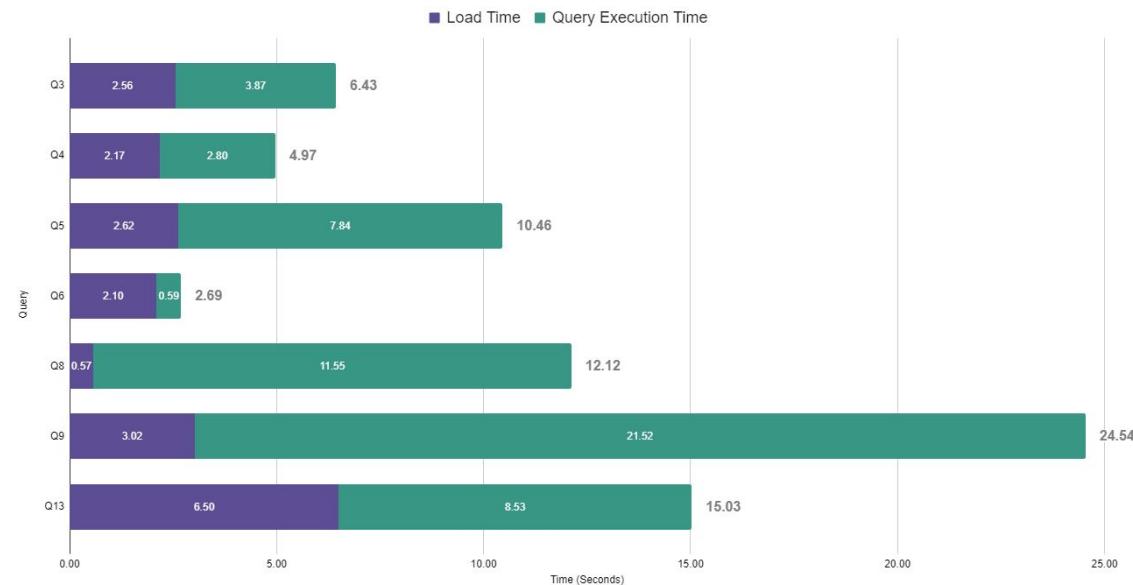
Streamz

Distributed stream  
processing using  
RAPIDS and Dask

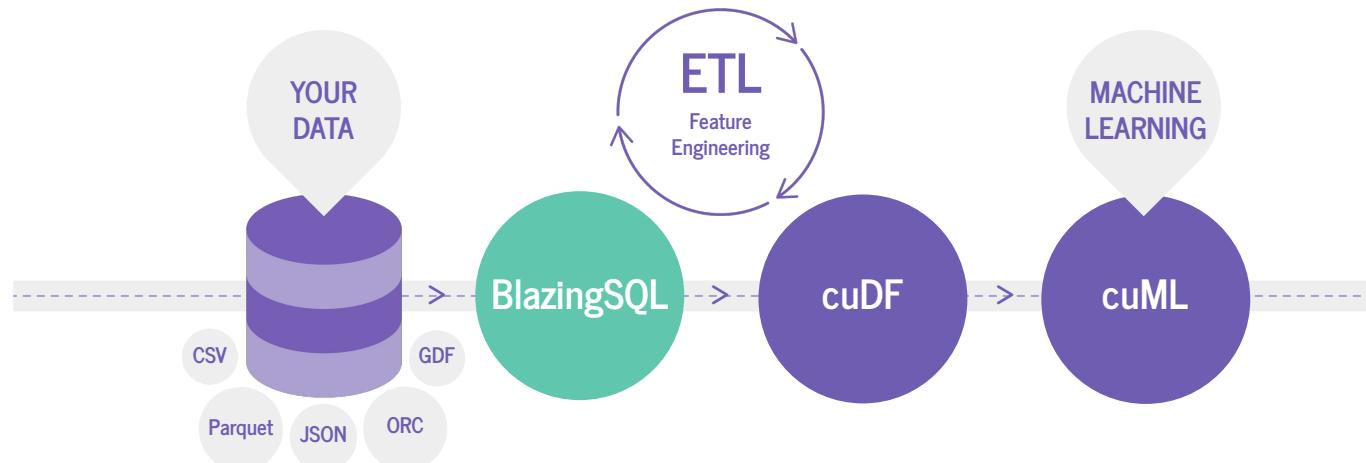
# BlazingSQL



TPC-H SF100 Query Times - NVME Storage



# BlazingSQL



```
from blazingsql import BlazingContext
import cudf

bc = BlazingContext()

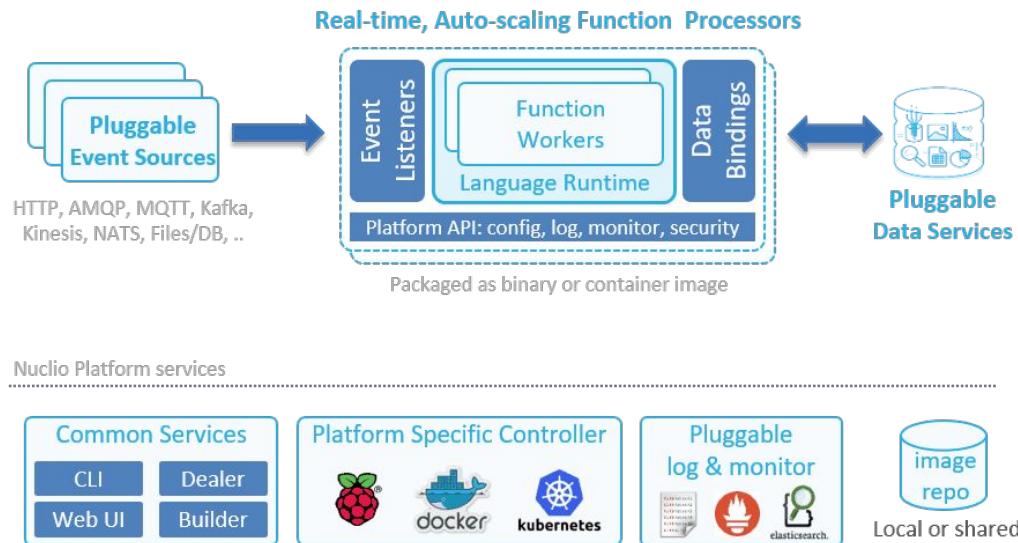
bc.s('bsql', bucket_name='bsql', access_key_id='<access_key>', secret_key='<secret_key>')

bc.create_table('orders', s3://bsql/orders/')

gdf = bc.sql('select * from orders').get()
```

# RAPIDS + Nuclio

Serverless meets GPUs



<https://towardsdatascience.com/python-pandas-at-extreme-performance-912912b1047c>

# Deploy RAPIDS Everywhere

Focused on robust functionality, deployment, and user experience



Google Cloud



Kubeflow



Azure



Azure Machine Learning



Cloud Dataproc



Amazon SageMaker



Alibaba Cloud

Integration with major cloud providers  
Both containers and cloud specific machine instances  
Support for Enterprise and HPC Orchestration Layers

# 5 Steps to getting started with RAPIDS

1. **Install RAPIDS** on using [Docker](#), [Conda](#), or [Colab](#)
2. **Explore** our [walk through videos](#), [blog content](#), our [github](#), the [tutorial notebooks](#), and our [examples workflows](#),
3. **Build** your own data science workflows.
4. **Join** our community conversations on [Slack](#), [Google](#), and [Twitter](#)
5. **Contribute** back. Don't forget to ask and answer questions on [Stack Overflow](#)

# Easy Installation

## Interactive Installation Guide

RAPIDS

HOME ABOUT GET STARTED COMMUNITY BLOG DOCS GITHUB

### RAPIDS RELEASE SELECTOR

RAPIDS is available as conda packages, docker images, and from source builds. Use the tool below to select your preferred method, packages, and environment to install RAPIDS. Certain combinations may not be possible and are dimmed automatically. Be sure you've met the required [prerequisites above](#) and see the [details blow](#).

METHOD	Conda	Docker + Examples	Docker + Dev Env	Source
RELEASE	Stable (0.9)		Nightly (0.10a)	
PACKAGES	cuDF	cuML	cuGraph	All Packages
LINUX	Ubuntu 16.04	Ubuntu 18.04		CentOS 7
PYTHON	Python 3.6		Python 3.7	
CUDA	CUDA 9.2		CUDA 10.0	
COMMAND	<pre>conda install -c rapidsai -c nvidia -c numba -c conda-forge -c anaconda \     cudf=0.9 cuml=0.9 cugraph=0.9 python=3.6 anaconda::cudatoolkit=9.2</pre>			
<a href="#">COPY COMMAND </a>				

# Explore: RAPIDS Github

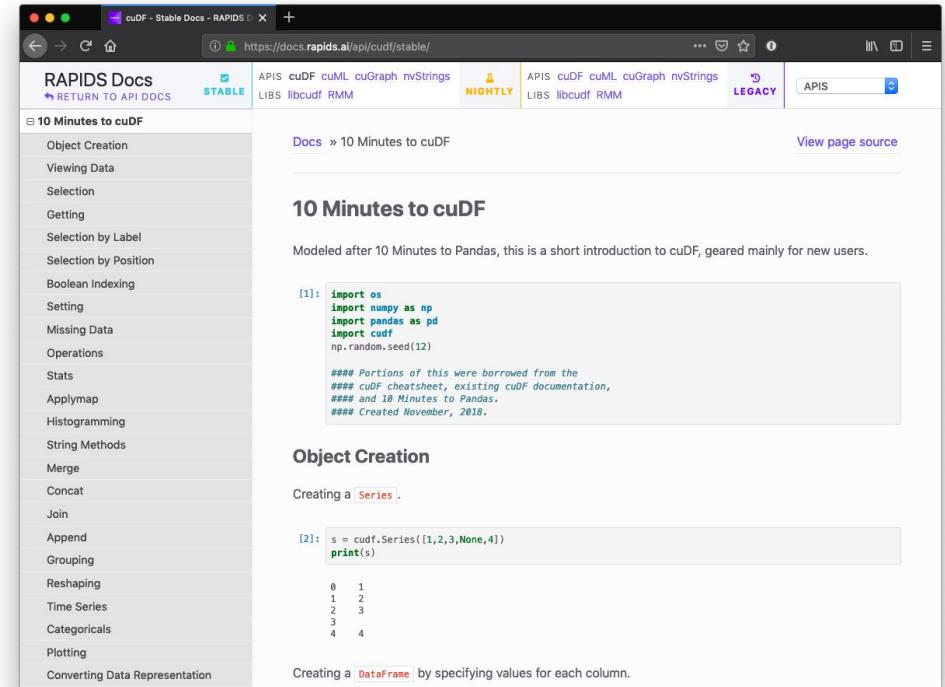
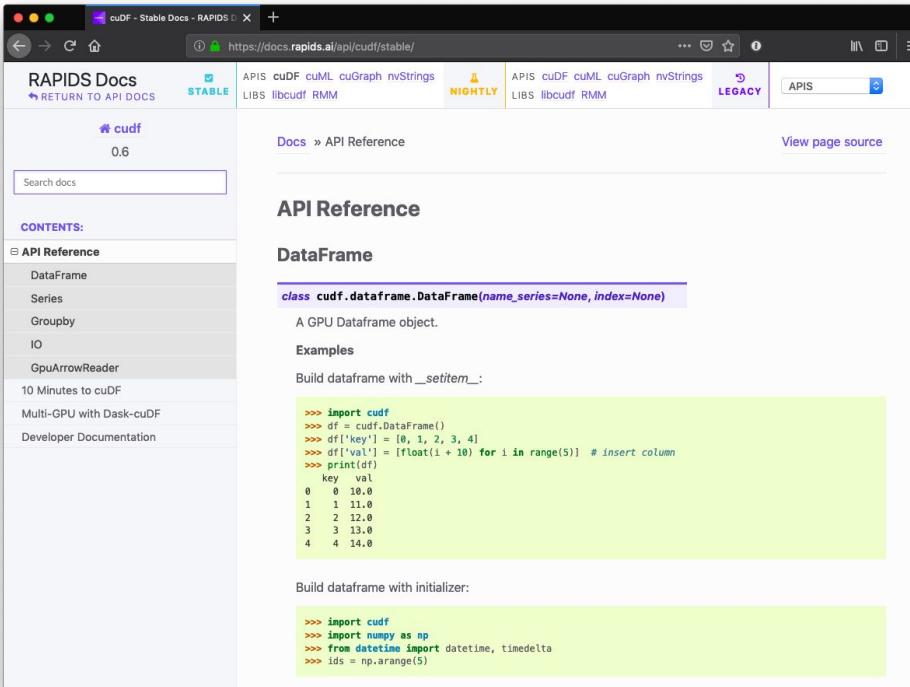
<https://github.com/rapidsai>

The screenshot shows the GitHub homepage for the RAPIDS repository. The header includes links for Pull requests, Issues, Marketplace, and Explore. Below the header is the RAPIDS logo and a brief description: "Open GPU Data Science". A link to the website (<http://rapids.ai>) is also present. The main navigation bar features tabs for Repositories (67), Packages, People (118), Teams (91), and Projects (6). The "Repositories" tab is currently selected. The page displays a section titled "Pinned repositories" containing six repository cards:

- cudf**: cuDF - GPU DataFrame Library. Cuda, 1.9k stars, 270 forks.
- cuml**: cuML - RAPIDS Machine Learning Library. C++ (indicated by a pink dot), 665 stars, 119 forks.
- cugraph**: cuGraph - RAPIDS Graph Analytics Library. Cuda, 204 stars, 52 forks.
- notebooks**: RAPIDS Sample Notebooks. Jupyter Notebook, 204 stars, 94 forks.
- notebooks-contrib**: RAPIDS Community Notebooks. Jupyter Notebook, 106 stars, 76 forks.
- cuxfilter**: GPU accelerated cross filtering. Python, 31 stars, 14 forks.

# Explore: RAPIDS Docs

Improved and easier to use!



<https://docs.rapids.ai>

# Explore: RAPIDS Code and Blogs

Check out our code and how we use it

README.md

## RAPIDS cuDF - GPU DataFrames

build running

NOTE: For the latest stable README.md ensure you are on the master branch.

Built based on the Apache Arrow columnar memory format, cuDF is a GPU DataFrame library for loading, joining, aggregating, filtering, and otherwise manipulating data.

cuDF provides a pandas-like API that will be familiar to data engineers & data scientists, so they can use it to easily accelerate their workflows without going into the details of CUDA programming.

For example, the following snippet downloads a CSV, then uses the GPU to parse it into rows and columns and run calculations:

```
import cudf, io, requests
from io import StringIO

url="https://github.com/plotly/datasets/raw/master/tips.csv"
content = requests.get(url).content.decode('utf-8')

tips_df = cudf.read_csv(StringIO(content))
tips_df['tip_percentage'] = tips_df['tip']/tips_df['total_bill']*100

# display average tip by dining party size
print(tips_df.groupby('size').tip_percentage.mean())

Output:
size
```



### RAPIDS Release 0.8: Same Community New Freedoms

Making more friends and building more bridges to more ecosystems. It's now easier than ever to get started with RAPIDS.



Josh Patterson

Jul 19 · 7 min read



### gQuant—GPU Accelerated examples for Quantitative Analyst Tasks

A simple trading strategy backtest for 5000 stocks using GPUs and getting 20X speedup



Yi Dong

Jul 16 · 6 min read ★



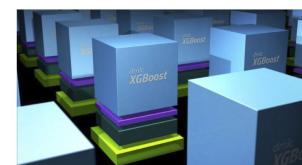
### Financial data modeling with RAPIDS.

See how RAPIDS was used to place 17th in the Banco Santander Kaggle Competition



Jiwei Liu

Jul 3 · 5 min read

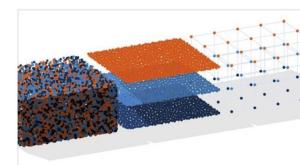


### NVIDIA GPUs and Apache Spark, One Step Closer

RAPIDS XGBoost4j-Spark Package Now Available

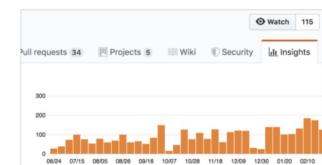


Karthikeyan Rajendran



### When Less is More: A brief story about XGBoost feature engineering

A glimpse into how a Data Scientist makes decisions about featuring engineering an XGBoost machine



### Nightly News: CI produces latest packages

Release code early and often. Stay current on latest features with our nightly conda and container releases.

<https://github.com/rapidsai>

<https://medium.com/rapids-ai>

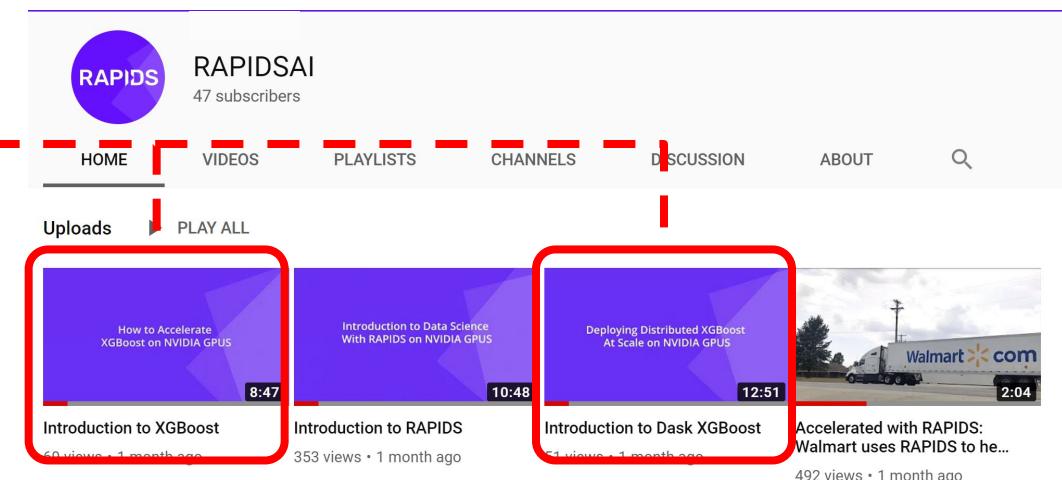
# Explore: Notebooks Contrib

Notebooks Contrib Repo has tutorials and examples, and various E2E demos. RAPIDS Youtube channel has explanations, code walkthroughs and use cases.

intro_tutorials	<a href="#">05_Introduction_to_Dask_cuDF</a>	This notebook shows how to work with cuDF DataFrames distributed across multiple GPUs using Dask.
intro_tutorials	<a href="#">06_Introduction_to_Supervised_Learning</a>	This notebook shows how to do GPU accelerated Supervised Learning in RAPIDS.
intro_tutorials	<a href="#">07_Introduction_to_XGBoost</a>	This notebook shows how to work with GPU accelerated XGBoost in RAPIDS.
intro_tutorials	<a href="#">08_Introduction_to_Dask_XGBoost</a>	This notebook shows how to work with Dask XGBoost in RAPIDS.
intro_tutorials	<a href="#">09_Introduction_to_Dimensionality_Reduction</a>	This notebook shows how to do GPU accelerated Dimensionality Reduction in RAPIDS.
intro_tutorials	<a href="#">10_Introduction_to_Clustering</a>	This notebook shows how to do GPU accelerated Clustering in RAPIDS.

## Intermediate Notebooks:

Folder	Notebook Title	Description
examples	<a href="#">DBSCAN_Demo_FULL</a>	This notebook shows how to use DBSCAN algorithm and its GPU accelerated implementation present in RAPIDS.
examples	<a href="#">Dask_with_cuDF_and_XGBoost</a>	In this notebook we show how to quickly setup Dask and train an XGBoost model using cuDF.



# Join the Conversation



[Google Groups](#)



[Docker Hub](#)



[Slack Channel](#)



[Stack Overflow](#)

# Contribute Back

Issues, feature requests, PRs, Blogs, Tutorials, Videos, QA...bring your best!

**cuml**  
cuML - RAPIDS Machine Learning Library  
machine-learning gpu machine-learning-algorithms  
cuda nvidia  
C++ Apache-2.0 111 608 186 (26 issues need help) 31 Updated 9 minutes ago

**cuDF**  
cuDF - GPU DataFrame Library  
anaconda gpu arrow machine-learning-algorithms  
h2o cuda pandas  
Cuda Apache-2.0 250 1,699 325 (6 issues need help) 41 Updated 31 minutes ago

**notebooks-contrib**  
RAPIDS Community Notebooks  
Jupyter Notebook Apache-2.0 56 70 10 (1 issue needs help) 8 Updated 40 minutes ago

**cugraph**  
cuGraph  
Cuda Apache-2.0 46 172 58 (1 issue needs help) 4 Updated 1 hour ago

TECH BLOG Walmart Labs ENGINEERING DATA SCIENCE INFOSEC UX DESIGN LEADERSHIP ABOUT

## How GPU Computing literally saved me at work?

Python+GPU = Power, 2 Days to 20 seconds

Abhishek Mungoli Follow May 9 · 9 min read

John Murray @MurrayData Follow

Comparison CPU vs GPU @rapidsai to project 100 million x,y points to lat/lon to 0.01mm accuracy. CPU 1 core c 65 mins, multicore c 13 mins, GPU #RAPIDS 2 seconds. I optimised the code since previous run. Dell T7910 Xeon E5-2640V4x2/NVIDIA Titan Xp cc @NvidiaAI @marc\_stampfli

```
John@Plato:~/Source/Python/misc$ python crs_test.py
Generating Data
CPU Iterative
4005.0377202 seconds
CPU mapped
3957.19386101 seconds
CPU multiprocessing
788.550751209 seconds
GPU Rapids
2.103230476 seconds
```

Getting Started with cuDF (RAPIDS)

Darren Ramsook Follow Jun 9 · 3 min read

# Getting Started

# RAPIDS Docs

## New, improved, and easier to use

The screenshot shows a web browser window for the cuDF - Stable Docs - RAPIDS Docs website at <https://docs.rapids.ai/api/cudf/stable/>. The page title is "cuDF - Stable Docs - RAPIDS Docs". The top navigation bar includes links for "APIS", "cuDF", "cuML", "cuGraph", "nvStrings", "LIBS", "libcudf", and "RMM". Below this, there are tabs for "NIGHTLY" and "LEGACY", with "APIS" currently selected. The main content area is titled "API Reference" under "DataFrame". It includes a class definition for `cudf.dataframe.DataFrame`, a description of it being a GPU Dataframe object, and examples of how to build a DataFrame using `__setitem__` and an initializer. The code examples are highlighted in green.

API Reference

### DataFrame

`class cudf.dataframe.DataFrame(name_series=None, index=None)`

A GPU Dataframe object.

**Examples**

Build dataframe with `__setitem__`:

```
>>> import cudf
>>> df = cudf.DataFrame()
>>> df['key'] = [0, 1, 2, 3, 4]
>>> df['val'] = [float(i + 10) for i in range(5)] # insert column
>>> print(df)
   key    val
0    0  10.0
1    1  11.0
2    2  12.0
3    3  13.0
4    4  14.0
```

Build dataframe with initializer:

```
>>> import cudf
>>> import numpy as np
>>> from datetime import datetime, timedelta
>>> ids = np.arange(5)
```

<https://docs.rapids.ai>

# RAPIDS Docs

## Easier than ever to get started with cuDF

The screenshot shows a web browser window titled "cuDF - Stable Docs - RAPIDS". The URL is <https://docs.rapids.ai/api/cudf/stable/>. The page content is the "10 Minutes to cuDF" guide. The left sidebar lists various topics under "10 Minutes to cuDF", including Object Creation, Viewing Data, Selection, Getting, Selection by Label, Selection by Position, Boolean Indexing, Setting, Missing Data, Operations, Stats, Applymap, Histogramming, String Methods, Merge, Concat, Join, Append, Grouping, Reshaping, Time Series, Categoricals, Plotting, and Converting Data Representation. The main content area starts with a header "10 Minutes to cuDF" and a note: "Modeled after 10 Minutes to Pandas, this is a short introduction to cuDF, geared mainly for new users." It includes a code block:

```
[1]: import os
import numpy as np
import pandas as pd
import cudf
np.random.seed(12)

#### Portions of this were borrowed from the
#### cuDF cheatsheet, existing cuDF documentation,
#### and 10 Minutes to Pandas.
#### Created November, 2018.
```

Below this, there is a section titled "Object Creation" with the sub-note: "Creating a `Series`." It shows another code block:

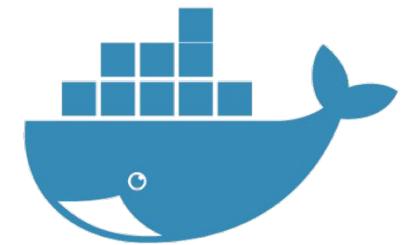
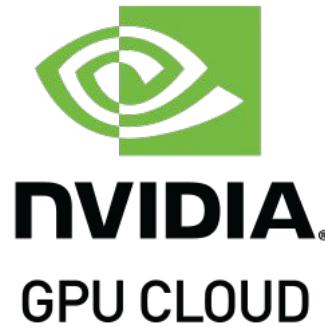
```
[2]: s = cudf.Series([1,2,3,None,4])
print(s)
```

0	1
1	2
2	3
3	
4	4

At the bottom, it says: "Creating a `DataFrame` by specifying values for each column."

# RAPIDS

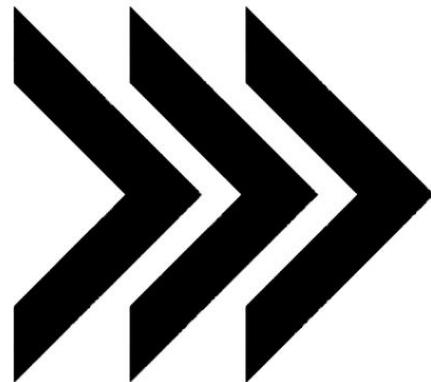
How do I get the software?



- <https://github.com/rapidsai>
- <https://anaconda.org/rapidsai/>
- <https://ngc.nvidia.com/registry/nvidia-rapidsai-rapidsai>
- <https://hub.docker.com/r/rapidsai/rapidsai/>

# Join the Movement

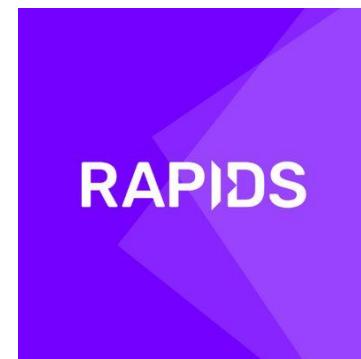
Everyone can help!



## APACHE ARROW

<https://arrow.apache.org/>

@ApacheArrow



## RAPIDS

<https://rapids.ai>

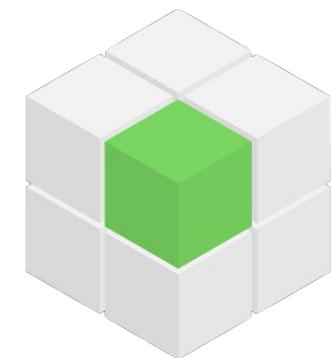
@RAPIDSAI



## Dask

<https://dask.org>

@Dask\_dev



## GPU Open Analytics Initiative

<http://gpuopenanalytics.com/>

@GPUOAI

Integrations, feedback, documentation support, pull requests, new issues, or code donations welcomed!

# THANK YOU

Joshua Patterson

joshuap@nvidia.com

@datametrician



# RAPIDS